4.7.5 The Open-Loop Excitation Collection

A second collection of 118 runs was recorded in the same deployment plant, with the same six-factor design and the same domain and split structure, and it is the one the adaptation module is trained on. It differs from the collection above in a single respect, and that respect is the point of it: there is no task controller. Each run is driven by open-loop excitation alone, the vehicle is stationary, and no policy is in the loop. The excitation is the same four-channel PRBS used as a probe above, at amplitude 0.02 and held for ten control steps, but it now drives the whole run rather than three windows inside one. Each run is 11 040 control steps, giving 1 302 720 steps across the collection.

The reason for collecting twice is structural. Under a task controller the command is computed from the state, so command and response are close to redundant and a plant descriptor has little left to explain; a model asked to identify the plant from that data is being given a question with no answer in it, and what it reports is a property of the data rather than of the architecture. Removing the controller makes the question well posed, which is why every model-level result in this thesis is fitted on this collection. How rich the excitation must be, and the amplitude beyond which plant identity stops being recoverable, remain outstanding measurements.

What each consumer takes from the recorded feature contract is set out in Section 4.8.2 for the adaptation module and Section 4.8.3 for the prediction models.

4.7.6 The Frozen Dataset

Both collections contain 118 complete runs under the identical domain design: 2 nominal, 40 hydraulic-only, 12 mechanical-only, 48 combined-training and 16 unseen combined-test. Collection ran in four vectorized IsaacLab environments with per-environment factors and logging, and slots that tripped a safety or validity check were discarded and retried individually. All 118 runs passed the final collection gate, which checks finite physics, exact model-grid alignment, command rank, factor-response authority, state coverage and saturation limits.

Each dataset is then *frozen*: the raw runs are checksummed and the run-level split is fixed at 82 training, 20 validation and 16 reserved combined-test runs, with zero factor-key leakage between them. The reserved sixteen do not influence history length, checkpoint, hyperparameter or architecture selection; they are scored only after each model is fixed, and test results are reported with validation results rather than in place of them.

4.7.7 Learning Timebase and Hydraulic State

The learning timebase is the actuator model grid, not the control clock. Samples are exact 50 ms ROM transitions taken at control indices $0, 3, 6, \ldots$: state row t and command row t predict state row $t + 3$. Nothing is interpolated. This is the same alignment rule as Section 3.5.1, and it is enforced rather than assumed, because a model trained across a mismatched interval is wrong by a fixed factor in the quantity it exists to supply.

The hydraulic state used by the system-dynamics surrogates is

$$\boldsymbol{s}_t = \left[\boldsymbol{x}_t^{1:4},\ \dot{\boldsymbol{x}}_t^{1:4},\ \Delta\boldsymbol{p}_t^{1:4}\right] \in \mathbb{R}^{12}, \tag{4.2}$$

with joint position and velocity carried as *context* only, and the prediction target the next 50 ms state delta. Differential pressure is not a measured signal. At runtime it is the actuator model's recursive state, generated by the command ROM or frozen forward model (Section 3.3). The system-dynamics surrogate and AAM may use it because it exists inside the deployed loop, but it inherits actuator-model error. Causal history is therefore retained to represent pressure transients that one sample cannot summarise.

4.8 Feature and Target Construction

4.8.1 The Actuator Models

The three learned actuator models take cylinder-space features, in the ordering and notation fixed in Section 5.1. Their input contracts differ because their questions differ. The inverse model receives the measured cylinder state, the logged rod forces, the desired increments, and the previously applied command, twenty-eight numbers in all. The forward model receives the applied valve command and the measured state, twenty numbers. The reduced-order model receives a six-sample history of state, previously predicted valve, and desired command, twenty-four numbers per step.

Three conventions apply throughout, and they are the same ones the models enforce at inference time.

Increment Targets

Every dynamic target is a within-run one-step increment, $\Delta \boldsymbol{x}_k = \boldsymbol{x}_{k+1} - \boldsymbol{x}_k$, never an absolute state. This removes the large static component of the signals and turns the trivial persistence solution into an explicit baseline rather than a hiding place.

Causal Alignment

Features at step k contain only quantities available at step k. The extraction stage verifies this by recomputing every logged target from the within-run difference of the states and requiring exact agreement, a check reported in Section 4.10.

Time-Base Conversion

The open-loop campaign is logged at 0.01 s but the models run at 0.05 s, so five consecutive raw samples are combined into one transition. The valve input attributed to that transition is the command held at the start of the window, which is the causally correct choice; the window mean is recorded for analysis only. Cylinder extension velocity is not logged directly in the closed-loop archive and is obtained by differentiating extension on the output grid, consistent with the trapezoidal identity that the strict data satisfies to $R^2 = 0.994$–0.9997 (Equation 5.6).

4.8.2 Features and Targets for the Adaptation Module

The adaptation module is trained on the open-loop collection of Section 4.7.5. Its history window carries only quantities the deployed stack produces at runtime: the five joint positions, the five joint velocities, the four cylinder pressure differences and the four

previously applied valve commands, giving the eighteen-value row of Section 5.5. The pressure entry is the learned actuator's own recursive state rather than a measurement (Section 3.3, Table 3.2); it is admitted because the model that generates it runs inside the deployed loop, and the consequence of that choice is examined where the estimator is evaluated. Four classes of signal are excluded by construction, and the exclusion is enforced by an explicit input mask rather than by convention:

- the physical factors $\mathbf{e}$ and the true latent $\mathbf{z}_{\mathrm{SD}}$, because recovering them is the estimator's job;
- any future observation, because the estimator is causal;
- the hidden internal state of the command ROM, because the deployment interface does not expose it;
- the desired task reference, so that the latent describes the plant rather than the trajectory being asked of it.

Windows are contained within a single run, and each carries the stopped-gradient target $\mathbf{z}_{\mathrm{SD}} = \mu(\mathbf{e})$ of its own run together with the run identifier, domain, split and valid history bounds. Applied to the 118 frozen runs this yields 444 624 causal history ticks; what is built from them belongs to Chapter 5.

The residual encoder deployed in AF shares this row but uses no collection at all, for the reason given in Section 4.8.4.

4.8.3 Features and Targets for the Prediction Models

The prediction models read the same runs through a third contract: a transition at the 50 ms tick rather than a history window, and the privileged factor vector $\mathbf{e}$, which the adaptation module never sees and which is what allows them to define the target the module is then asked to recover. That contract is set out with the dataset it produces in Section 4.11, because what it is for is only visible alongside the audit it makes possible.

4.8.4 On-Policy Data for the Deployed Adaptation Residual

The residual that carries the adaptation result in AF is not trained from the frozen collection above. It has no dataset, no stopped-gradient latent target and no run-level splits, because its encoder never regresses a target: the encoder and the residual head are optimised together under the reinforcement objective, so their training data is the on-policy experience the stack generates as it learns.

What replaces the dataset is a plant design, since on-policy training makes the training distribution a property of the environment rather than of a file. The cutter mass is drawn from two well-separated modes and flips between them *mid-episode*, so no single compromise response is correct in both halves of a run. The flip instant is jittered per environment so it cannot be anticipated on a timer, and the direction is balanced across environments so it cannot be assumed. History is therefore not merely available to the

residual but necessary to it, and an arm that cannot see the history falls well below the frozen base rather than merely failing to improve on it (Section 7.5).

Two properties of the frozen collection are inherited unchanged. The history row is the same eighteen-value row defined above, with the same four exclusion classes enforced by the same input mask, and the plant is still the deployment plant, so the change of training regime does not substitute a different system.

4.9 Preprocessing and Data Splits

Splitting is done at the level of complete runs, never of individual samples, because a random sample-level split would place transitions from the same trajectory on both sides of the boundary and report a score that mostly measures interpolation within a run. For the command-conditioned dataset the 1 040 runs are divided into 936 training, 52 validation and 52 test runs, giving 407 160, 22 620 and 22 620 transitions, with no run identifier in more than one split. The IsaacLab deployment collection is split by run *and* by factor configuration (82/20/16), the stronger condition: separating runs alone would let two runs of the same plant fall on opposite sides of the boundary and turn a claim about unseen plants into a claim about unseen episodes.

The strict open-loop export needs different treatment because it contains only four runs, and holding one out would remove an entire plant draw from training. Each run is instead divided into twenty temporal blocks, every fifth block is assigned to validation, and a five-sample guard interval is discarded at each boundary so that no 0.05 s transition window spans the split. This preserves comparable target distributions in training and validation.

Normalisation statistics are computed on the training split only and stored inside each checkpoint, so a model always normalises with the statistics it was fitted with; the loader enforces the stored contract rather than recomputing it. Samples are carried with an explicit validity flag rather than silently dropped, 93.7 % of the strict export being marked valid and the remainder being the excluded startup and shutdown seconds and the zero-command gaps. Keeping them with a flag means the time base stays contiguous and the within-run difference checks remain meaningful.

4.10 Dataset Diagnostics

No dataset in this thesis is used before it passes an explicit audit, for the reason given at the end of Section 4.4. The audit of the strict open-loop export is reproduced in Table B.4.

The two rank figures speak directly to requirement R2. The instantaneous command matrix has rank four, so no channel is a linear combination of the others, and the lagged regressor matrix built from the command history has rank eighty, so the excitation is persistently exciting to the order the models need. The condition number of that lagged matrix is about 4.6×10^3, finite and workable rather than comfortable, and it is the price paid for the stroke coverage the long plateaus buy.

The three increment checks return a maximum error of exactly zero, which is expected rather than suspicious: the targets are computed from the logged states by the same within-run difference the check recomputes, so the check verifies that the extraction stage did what it claims rather than that the simulation is accurate. It catches misalignment,

off-by-one indexing and rows reordered or concatenated across run boundaries. Physical fidelity is addressed separately by the validation study of Section 3.4.4.

Two properties of the Simscape material are recorded here because later chapters refer to them. Stroke coverage is 12 % on the head-pivot cylinder, the narrowest of the four channels and also the one the strict forward model predicts least well. The rod-force channel carries zero supervised weight in the reduced-order model, a design choice rather than a data defect: force is generated internally by the PhysX plant at deployment, and the logged channel additionally fails a piston-area consistency check on this export. No result rests on it.

The IsaacLab collection is audited against different failure modes: design adequacy before use (unique factor keys, full active-factor rank, disjoint keys across splits), physical validity per run (finite physics, exact model-grid alignment, full command rank, bounded saturation), and leakage structurally, with privileged channels removed by an explicit input mask rather than by trusting the feature list. Its limitation is scale, 118 runs over six factors being a sparse design, which is why the paired test of Section 4.11.1 is reported with bootstrap intervals and per-run detail rather than as a pooled point estimate. That limitation does not bound the closed-loop adaptation result, which is estimated from paired environments generated on-policy under the plant design of Section 4.8.4.

4.11 The System-Dynamics Dataset

The prediction models read the deployment collection through a third contract, built once and frozen with the rest. It is a derived view rather than a separate collection: the same 118 runs, resampled from the 60 Hz control log to the 20 Hz model tick by taking every third step, giving 3 840 transitions per run and 453 120 in total. The run-level split is inherited unchanged, 82 training, 20 validation and 16 test runs, so a run contributes to exactly one split here and in every other view of this data.

Each sample is one transition. The input is the 12-dimensional hydraulic state, the 10-dimensional joint context and the valve command pair, previous and current, for 30 channels; the target is the hydraulic-state increment to the next tick in the same 12 dimensions. Training draws five consecutive transitions at a time, so the objective is a short rollout rather than a one-step fit. Four per-tick flags record cylinder limit, command saturation, valve saturation and non-finite physics, so a model can be scored on valid ticks alone. Each run carries its own six-factor vector $\mathbf{e}$, which is privileged: it is what allows a target to be defined for the adaptation module of Section 4.8.2, and it is masked out of every deployable input.

4.11.1 Dataset Separation Across Plants

One property of this dataset cannot be checked by an audit of its structure. The checks of Section 4.10 establish that the runs are valid, that the factor design is full rank and that the splits do not leak, but not that the runs *differ by plant* in a way any model could use. If they do not, nothing built on this dataset can recover the difference.

The test is a paired one on the dataset itself. Two prediction models are trained on identical samples, splits and objective; one receives a two-dimensional plant context from

an encoder $\mathbf{z}_{\mathrm{SD}} = \mu(\mathbf{e})$ trained jointly with it, and the matched control receives no such input. Neither is an actuator and neither enters a control stack.

The context helps, and narrowly. Pooled rollout error falls by 7.20×10^5 on the 20 validation runs, with a 95 % interval of $[-1.03 \times 10^6, -4.30 \times 10^5]$, and by 4.41×10^5 on the 16 test runs, interval $[-7.61 \times 10^5, -1.74 \times 10^5]$. Both exclude zero, so the dataset does separate plants. Where it separates them is the more useful result: the per-channel decomposition puts the whole gain in pressure, and the extension and velocity intervals cross zero on both pools.

Table 4.1: Per-channel decomposition of the paired advantage at a 20-tick horizon, pooled over 36 held-out runs. Values are the mean paired difference between the context-conditioned model and the matched control, with a 95 % bootstrap interval. The advantage is confined to pressure.

Channel	mean paired difference	95 % interval	excludes zero
extension	5.2×10^{-5}	$[-3.6 \times 10^{-5},\ 1.4 \times 10^{-4}]$	no
velocity	-2.4×10^{-4}	$[-7.1 \times 10^{-4},\ 2.4 \times 10^{-4}]$	no
pressure	1.75×10^5	$[2.1 \times 10^4,\ 3.3 \times 10^5]$	yes

Two limits travel with this result and both matter later. It compares separately trained models, so it establishes that the context is informative, not that an estimator could recover it or that a controller could use it; Section 5.6.3 substitutes contexts within one model, which is the tighter test. And the informative channel is pressure, which the controller does not integrate. That is the first appearance of the delivery problem the rest of this thesis runs into: this dataset carries plant identity in the one channel that reaches the command path least directly.

With the data described and audited, the next chapter turns to what was learned from it: the inverse and forward models that form the IF actuator path, the command-conditioned reduced-order model that replaces them in R, and the adaptation module that estimates at runtime what the surrogate above was given.

5. Learned Models and Online Adaptation

Chapter 4 established how the datasets were generated and why both closed-loop operation and open-loop valve excitation were required. This chapter develops the learned representations that consume them: the inverse-dynamics (ID) and forward-dynamics (FD) networks of the IF actuator path, the command-conditioned reduced-order model (ROM) that replaces them in R, and an offline Arm Adaptation Module (AAM) study that tests whether plant context can be recovered from causal history. This offline AAM and its ID adapter are diagnostic prototypes. The deployed AF encoder and residual are trained jointly under PPO in Chapter 6.

The chapter continues one evidence chain. Section 4.11 established the first link on the data itself: a compact plant context exists and improves prediction, and the improvement sits in pressure. What follows is whether that context is recoverable causally from observable history, whether the recovered estimate is as informative as a privileged one, and whether it reaches the command. The chain fails at that last link, and the diagnosis specifies the redesign carried out in Chapter 6.

One commitment organises the chapter: no model is accepted on a supervised score, because such scores repeatedly proved uninformative here, a forward model reaching a one-step R^2 near 0.999 while having learned almost nothing about valve authority (Section 5.3). Every model is therefore presented with its contract, training objective and the offline evidence that accepted or refused it. All models operate on cylinder-space signals.

5.1 Modelling Overview and Common Notation

All learned models act on the four hydraulic cylinders of the MULAG arm, indexed $i \in \{1, \dots, 4\}$ in the fixed order slew cylinder (*Drehzapfen*), boom cylinder I (*Ausleger I*), boom cylinder II (*Ausleger II*) and head pivot cylinder (*Messerkopf-Schwenk*). The passive fifth joint is not actuated and enters no model input. At the learned-model sample time $T = 0.05\,\mathrm{s}$ the discrete index k refers to the cylinder extension $\boldsymbol{x}_k$, the extension velocity $\dot{\boldsymbol{x}}_k$, the chamber pressure difference $\Delta\boldsymbol{p}_k$ and the logged rod force $\boldsymbol{F}_k$, each in $\mathbb{R}^4$. The

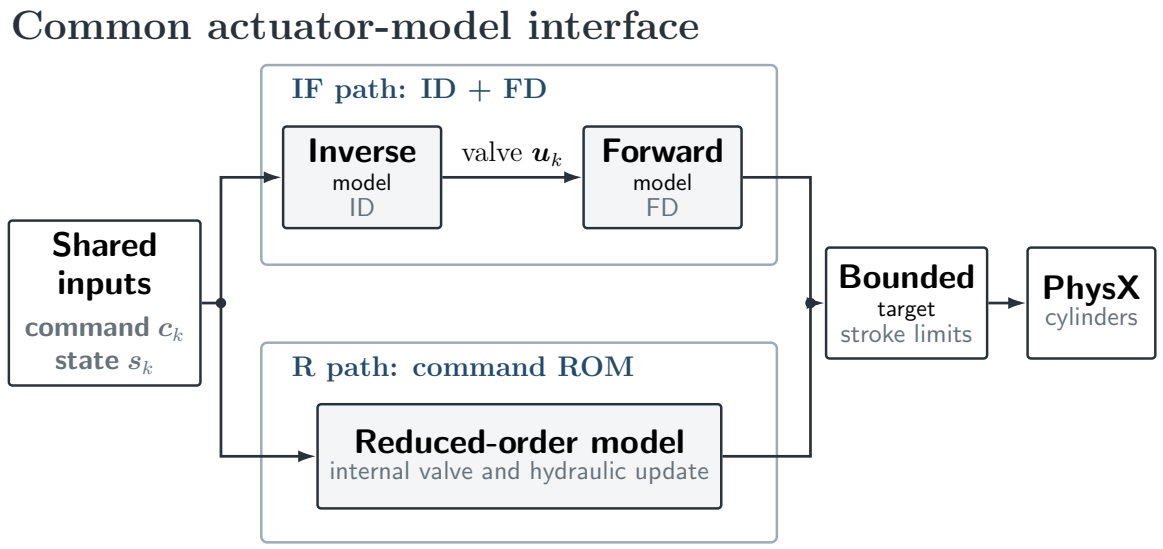

Figure 5.1: Common interface of the two alternative actuator paths. Both receive the desired cylinder command and measured hydraulic state, and both produce the same bounded PhysX cylinder target. The IF path composes the inverse and forward models; the R path uses one command-conditioned reduced-order model.

valve command $\boldsymbol{u}_k \in [-1, 1]^4$ is the normalised spool command of the four proportional valves. The desired cylinder command produced by the task controller is the pair

$$\boldsymbol{c}_k = \left[\Delta \boldsymbol{x}_k^*,\ \Delta \dot{\boldsymbol{x}}_k^*\right] \in \mathbb{R}^8, \tag{5.1}$$

the extension and extension-velocity increments to be realised over the next sample. Predictions carry a hat, and $\Delta(\cdot)_k = (\cdot)_{k+1} - (\cdot)_k$ denotes a within-run one-step increment throughout.

Figure 5.1 compares the actuator paths. In IF, the ID maps the state and desired increment to a valve command, and the FD predicts the resulting increment. In R, one command-conditioned model predicts the valve internally and propagates pressure, velocity and extension. Both paths produce the same bounded PhysX cylinder target in IsaacLab [Mittal 23] under the same policy action contract [Schulman 17].

Three conventions apply throughout. Dynamic targets are within-run increments, with $\Delta\hat{\boldsymbol{x}} = \mathbf{0}$ as the persistence baseline. Each checkpoint uses its own training-set normalisation. Inputs at step k are causal, and targets are measured transitions to step $k + 1$.

Simscape supplies the identification data described in Chapter 4. The exported arrays are audited, frozen and used to train the models in PyTorch [Paszke 19]; Simscape is not part of training or deployment.

The closed-loop Simscape export [The MathWorks, Inc. 24a] contains 452 400 transitions at $T = 0.05\,\mathrm{s}$ from 1 040 runs, split by run into 936/52/52 training, validation and test runs. It supplies the supervised ID, FD and command-ROM data. The FD repair uses four open-loop 64-s runs at 0.01 s with rank-4 valve excitation on four plant draws.

5.2 Inverse Dynamics Model

The ID network maps the current hydraulic state and desired cylinder increment to a valve-command change:

$$\Delta \boldsymbol{u}_k = f_{\mathrm{ID}}\left(\boldsymbol{x}_k,\, \dot{\boldsymbol{x}}_k,\, \Delta \boldsymbol{p}_k,\, \boldsymbol{F}_k,\, \Delta \boldsymbol{x}_k^*,\, \Delta \dot{\boldsymbol{x}}_k^*,\, \boldsymbol{u}_{k-1}\right) \in \mathbb{R}^4, \tag{5.2}$$

a 28-dimensional input. The deployed command is $\boldsymbol{u}_k = \mathrm{clip}(\boldsymbol{u}_{k-1} + \Delta \boldsymbol{u}_k,\, -\boldsymbol{u}_{\max},\, \boldsymbol{u}_{\max})$ with per-axis limits $\boldsymbol{u}_{\max} = [0.95,\, 0.95,\, 0.8,\, 0.8]$ stored in the checkpoint. The change formulation retains the previous valve operating point near the direction-dependent dead zone (Section 2.1.2). Logged achieved transitions serve as desired increments because no independent cylinder-space reference was recorded.

The 512-wide residual MLP [He 16] uses SiLU activations, layer normalisation and dropout [Elfwing 18, Ba 16, Srivastava 14]. Adam minimises a smooth-L_1 loss on normalised targets, with saturated samples up-weighted [Kingma 15, Huber 64].

Figure 5.2 separates the three input groups and shows the four residual blocks explicitly. The network predicts a four-axis valve-command increment, which is accumulated and clipped before use.

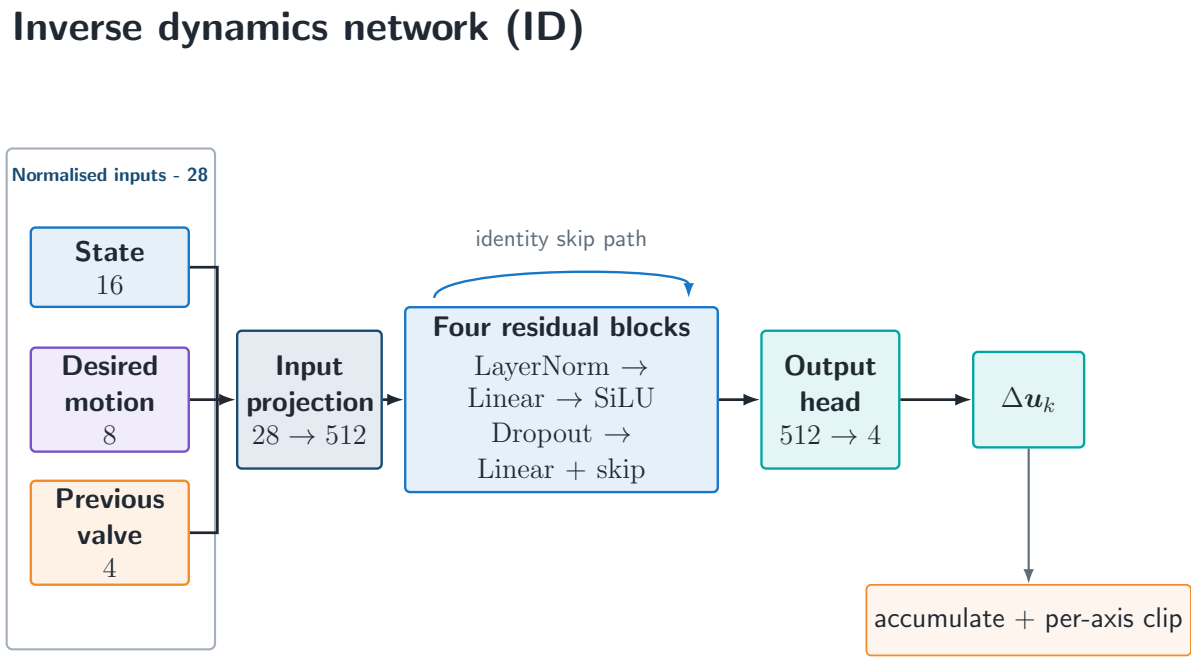

Figure 5.2: Inverse-model architecture. Current hydraulic state (16), desired motion (8) and the previous valve command (4) form the 28 inputs. A 512-wide, four-block residual MLP predicts $\Delta \boldsymbol{u}_k$, which is accumulated and clipped per axis to form the bounded valve command.

Per-valve validation R^2 is 0.9973–0.9996. This shows reproduction of logged commands, not plant authority: in closed-loop data, achieved transitions and logged commands are correlated (Section 2.4.2).

5.2.1 Paired Control Fine-Tuning

The composite is assessed by the *paired differential response*: the FD prediction under the ID command for the requested target minus the prediction under its command for a zero target,

$$\begin{aligned} \boldsymbol{u}_k^{\mathrm{tgt}} &= f_{\mathrm{ID}}(\cdot,\, \boldsymbol{c}_k,\, \boldsymbol{u}_{k-1}), \qquad \boldsymbol{u}_k^{0} = f_{\mathrm{ID}}(\cdot,\, \mathbf{0},\, \boldsymbol{u}_{k-1}), \\ \Delta\hat{\boldsymbol{x}}_k^{\mathrm{pair}} &= f_{\mathrm{FD}}\Big(\boldsymbol{u}_k^{\mathrm{tgt}}, \cdot\Big) - f_{\mathrm{FD}}\Big(\boldsymbol{u}_k^{0}, \cdot\Big). \end{aligned} \tag{5.3}$$

Before fine-tuning, the composite normalised control error was 1.00, equal to the zero-response baseline, and all regression slopes were below 0.004 (Table 5.1). The networks were fine-tuned jointly on $\mathcal{L}_{\mathrm{ctl}} = \|\Delta\hat{\boldsymbol{x}}^{\mathrm{pair}} - \Delta\boldsymbol{x}^{*}\|^2_{\mathrm{norm}} + \lambda_{\mathrm{ID}}\mathcal{L}_{\mathrm{ID}} + \lambda_{\mathrm{FD}}\mathcal{L}_{\mathrm{FD}}$, with the supervised terms retained. Target blocks were shuffled across valid states to break the logged state–target pairing. After 30 epochs, control error fell to 0.109; all slopes exceeded 0.8 and all correlations exceeded 0.92. This frozen checkpoint is the inverse model of the IF actuator path, and of the latent-conditioned adapter route of Section 5.6. It is not part of AF, whose command path contains no inverse model (Section 6.7.1).

Table 5.1: Effect of paired control fine-tuning on the composite differential response (5.3), per cylinder axis on validation data. Before fine-tuning the composite had essentially no control authority despite near-perfect supervised scores of both networks; the composite normalised control error fell from 1.0011 to 0.1086.

	Regression slope		Correlation	
Cylinder axis	before	after	before	after
1 (slew)	0.003	0.818	0.135	0.920
2 (boom I)	0.002	0.928	0.090	0.967
3 (boom II)	0.001	0.849	0.037	0.933
4 (head pivot)	0.004	0.912	0.249	0.958

5.2.2 Transfer of the Frozen Inverse Model

Because the adapter corrects a frozen ID, its transfer was measured on all 118 frozen runs of Section 4.7, comprising 2 nominal, 40 hydraulic-only, 12 mechanical-only and 64 combined runs, of which 16 are held-out combined-test. These runs lie outside the ID's Simscape-only training distribution and are recorded in the plant the adaptive controller is closed around. Two regimes were scored: a teacher-forced one-step transfer supplying the measured cylinder state at every tick, and a closed-loop transfer feeding back the pair's own prediction while supplying the recorded reference. Pressure and force are held at their training means, as the deployed feature contract requires.

Teacher-forced, the pair transfers well and evenly (Table 5.2). One-step extension error is millimetre-scale in every domain and nearly flat across them, from 2.36 mm on the combined runs to 3.36 mm on nominal, and the held-out combined-test group is no worse than the domains used in training, at 2.34 mm. This is the regime in which IF uses the pair, because the policy closes the loop at every tick and never asks it to predict its own future.

Table 5.2: Transfer of the frozen inverse and forward pair across deployment domains, all 118 runs of the deployment collection. The one-step columns are teacher-forced; the saturation column is the fraction of pre-clip valve elements at the limit during closed-loop transfer. The ROM valve recorded on these same runs never saturates.

		one-step Δx error		closed-loop
Domain	runs	RMSE	worst run	valve saturation
combined	64	2.36 mm	5.72 mm	0.888
hydraulic	40	2.43 mm	4.96 mm	0.885
mechanical	12	2.81 mm	4.14 mm	0.919
nominal	2	3.36 mm	4.02 mm	0.929
all	118	—	—	0.891

Closed-loop, it does not transfer at all, and the failure is uniform rather than occasional. Every one of the 118 runs leaves the recorded reference by more than one metre, after about 32 s on the combined and hydraulic domains and sooner on the harder ones. The mechanism is saturation: the pair's pre-clip valve command sits at its limit on 89.1 % of elements over the collection, between 88.5 and 92.9 % by domain, while the ROM valve recorded on those same runs never reaches its bound at all. For scale, the measured plant tracks the recorded cylinder reference at about 0.807 mm on the held-out group, so the closed-loop divergence is a failure of model transfer and not a property of the collection.

Two consequences follow. The pair cannot be promoted as an unadapted autoregressive deployment path, because its own predicted state leaves the support its training data covered, though nothing in that result disturbs the one-step use IF makes of it. What it does disturb is adaptation: a bounded correction added to a command already resting on its stops has nothing left to move (Section 6.7.1).

5.3 Forward Dynamics Model

The FD network predicts the one-step cylinder response to a valve command,

$$\left[\Delta\hat{\boldsymbol{x}}_k,\ \Delta\hat{\dot{\boldsymbol{x}}}_k,\ \Delta\widehat{\Delta\boldsymbol{p}}_k\right] = f_{\mathrm{FD}}\left(\boldsymbol{u}_k,\ \boldsymbol{x}_k,\ \dot{\boldsymbol{x}}_k,\ \Delta\boldsymbol{p}_k,\ \boldsymbol{F}_k\right) \in \mathbb{R}^{12}, \tag{5.4}$$

using the same residual trunk as the ID network. Supervised pretraining on the closed-loop export, with an auxiliary recursive rollout loss [Nagabandi 18], reached grouped one-step R^2 near 0.999 for the extension increments.

Figure 5.3 shows the forward model separately because its input and output contracts differ from the inverse model despite the shared residual trunk.

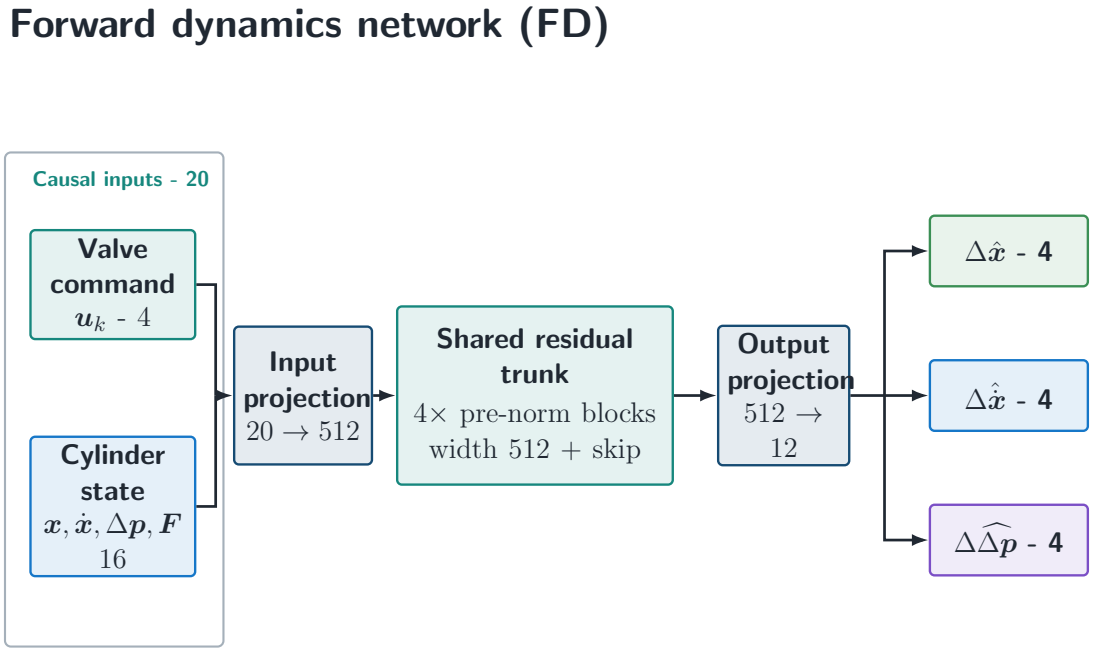

Figure 5.3: Forward-model architecture. The applied valve command and cylinder state condition a one-step hydraulic transition through a 512-wide, four-block residual MLP. Its 12 outputs are extension, velocity and pressure increments. Extension is the integrated controller output; velocity and pressure are diagnostic outputs.

5.3.1 Assessment of the Supervised Score

The supervised score did not establish valve sensitivity. At a 0.5 s horizon, zeroing the valve input changed the rollout-error ratio to only 0.94. The FD had mainly extrapolated the smooth closed-loop kinematics because command and state history were nearly redundant. Closed-loop ablation confirmed that the valve was not identifiable, so repair required new data.

The repair dataset applies rank-4 open-loop valve plateaus to four plant draws. Its trajectories are distinct, valve history predicts pressure with $R^2 \approx 0.80$–0.85, valve-regressor variance-inflation factors are 7–10, and relief saturation occupies 9–22 % of samples.

Starting from the paired-control checkpoint, strict retraining minimised

$$\mathcal{L}_{\mathrm{FD}} = \mathcal{L}^{\mathrm{strict}}_{\mathrm{Huber}} + 0.25\,\mathcal{L}^{\mathrm{MAT}}_{\mathrm{Huber}} + 2.0\,\mathcal{L}^{(10)}_{\mathrm{roll}} + 10.0\,\mathcal{L}_{\mathrm{kin}}, \tag{5.5}$$

combining the one-step loss on the strict data, a replay term on the original closed-loop data, a 0.5 s recursive rollout loss, and a kinematic consistency term enforcing the trapezoidal identity

$$\Delta \boldsymbol{x}_k = T\,\dot{\boldsymbol{x}}_k + \tfrac{1}{2}\,T\,\Delta\dot{\boldsymbol{x}}_k, \tag{5.6}$$

which holds in the strict data with $R^2 = 0.994$–0.9997. Valve authority varies by a factor of 15–75 across pressure quartiles and can change sign, so the ROM represents this dependence by operating point [Merritt 67].

Strict retraining reduces extension rollout RMSE from 75.2 mm to 12.4 mm at 0.5 s and from 157.9 mm to 27.3 mm at 1.0 s. The valve-zero error ratio rises from 0.94 to 1.81, and per-axis one-step extension R^2 improves from $[-5.7, 0.60]$ to $[0.76, 0.98]$. Zero-order hold remains competitive below 0.25 s, and velocity increments remain weak on two axes. Deployment therefore integrates only $\Delta\hat{\boldsymbol{x}}$ and retains velocity as a diagnostic output.

5.4 Reduced-Order Model

The R path uses one command-conditioned model instead of the ID–FD composition. Its structure is grey-box: physical equations propagate pressure, force, velocity and extension. Neural heads estimate valve realisation, load and bounded residuals [Karniadakis 21, Lutter 19]. Valve authority is pressure dependent by construction through the orifice-flow and continuity equations [Merritt 67]. A GRU [Cho 14] reads six samples (0.25 s) of cylinder state, previous predicted valve command and desired command. The measured valve is only an internal-head training target [Hwangbo 19].

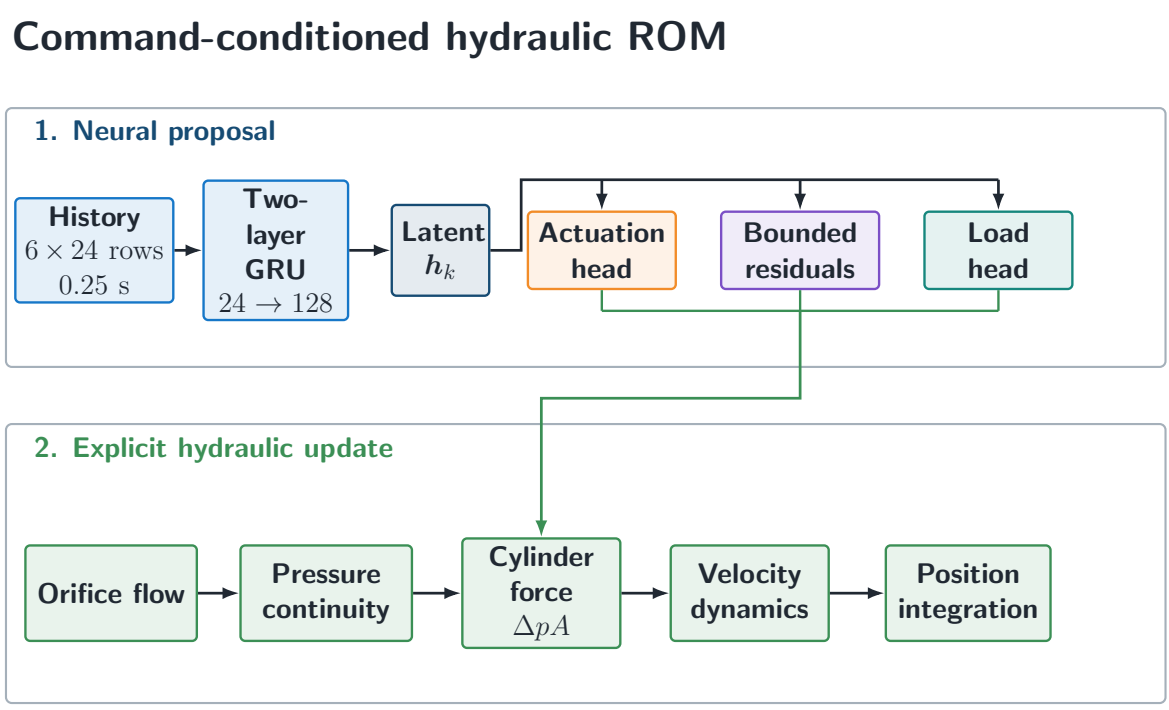

Figure 5.4: Command-conditioned ROM. A recurrent neural proposal encodes a six-step history. Its actuation, bounded-residual and load heads condition the explicit hydraulic update: orifice flow, pressure, force, velocity and position. The explicit chain constrains the predicted transition.

The actuation head uses a learnable smooth dead zone with per-cylinder width and asymmetric gains [Yang 26]. Its output $\hat{a}_i \in [-1, 1]$ drives the orifice-flow and continuity

update

$$q_i = \hat{a}_i\, k_{q,i} \sqrt{\tfrac{1}{2}\big(p_{s,i} - \operatorname{sgn}(\hat{a}_i)\,\Delta p_i\big)_+},$$
$$\Delta\hat{p}_{i,k+1} - \Delta p_{i,k} = \frac{4\beta_i}{V_i(x_i)}\big(q_i - A_i\dot{x}_i - k_{\ell,i}\,\Delta p_i\big)\,T \;+\; r_{p,i}, \tag{5.7}$$

where β_i, $p_{s,i}$, $k_{q,i}$, $k_{\ell,i}$, A_i and $V_i(x_i)$ are the effective bulk modulus, supply pressure, flow gain, leakage, piston area and chamber volume. Pressure then produces force and motion:

$$\hat{F}_i = \Delta\hat{p}_{i,k+1}A_i + r_{F,i}, \qquad \ddot{x}_i = \frac{\hat{F}_i - c_{v,i}\,\dot{x}_i - \ell_i(\boldsymbol{h}_k)}{m_i},$$
$$\dot{\hat{x}}_{i,k+1} = \dot{x}_{i,k} + \operatorname{sat}\big(T\,\ddot{x}_i + r_{v,i}\big), \qquad \hat{x}_{i,k+1} = x_{i,k} + \tfrac{T}{2}\big(\dot{x}_{i,k} + \dot{\hat{x}}_{i,k+1}\big) + r_{x,i}, \tag{5.8}$$

Here $c_{v,i}$ is viscous friction, m_i is effective modal mass and $\ell_i(\boldsymbol{h}_k)$ is the learned configuration-dependent load. Physical parameters are trained in log-space inside fixed clamps. Several settle on a clamp bound and are therefore fitting coefficients rather than measured constants: modal mass and viscous friction both reach their upper clamp of 10^7 (in kilograms and in newton-seconds per metre respectively), the piston area reaches its lower clamp of $1\,\mathrm{cm}^2$ against a $64\,\mathrm{cm}^2$ initialisation, and the supply pressure separates onto both clamps, with two cylinders at 1 MPa and two at 50 MPa. The same values appear in every checkpoint of this family. Pressure and velocity residuals are bounded to 10 % of their step limits, and the position residual to 1 cm, capping the magnitude each head may add per step.

The explicit chain and the neural heads contribute to the transition in very different proportions. With the residual heads inactive the explicit chain alone produces no net travel, its velocity increments falling six orders of magnitude below their nominal values: at the fitted parameters the friction term $c_{v,i}\dot{x}_i$ reaches 1.0 MN at 0.1 m/s, against a pressure force $\Delta\hat{p}_i A_i$ of about 10 kN at 100 MPa. Realised displacement is produced by the residual heads. The pressure equation nonetheless remains active and responsive to its own parameters: with the heads inactive, halving $k_{\ell,i}$ still moves Δp by 63 % of its range. Hydraulic parameters therefore reach the predicted motion indirectly, through the pressure state entering the recurrent encoder, rather than through the force balance. The model has a grey-box structure. Its equations fix the transition form and keep it differentiable in named hydraulic quantities, but they do not identify the machine's physical constants.

Training uses rollout horizons $H \in \{1, 5, 20, 50\}$ and scheduled sampling up to 75 % self-fed [Bengio 15]. Residual penalties retain the physical structure. Logged rod force has zero supervised weight because PhysX generates it at deployment and the export fails the piston-area consistency check.

With the true command, mean one-step R^2 is 0.998 for the internal valve, 0.64 for velocity and 0.97 for extension. Zeroing or shuffling the command reduces velocity R^2 below zero and extension to about 0.5; the extension Jacobian is diagonally dominant with condition number 2.14. At 5 s, terminal extension R^2 is 0.998, 0.998, 0.84 and 0.91 across the cylinders. Pressure-increment R^2 is only 0.06–0.09, so the internal pressure is a dynamic state, not a calibrated pressure estimate (Figure 5.5).

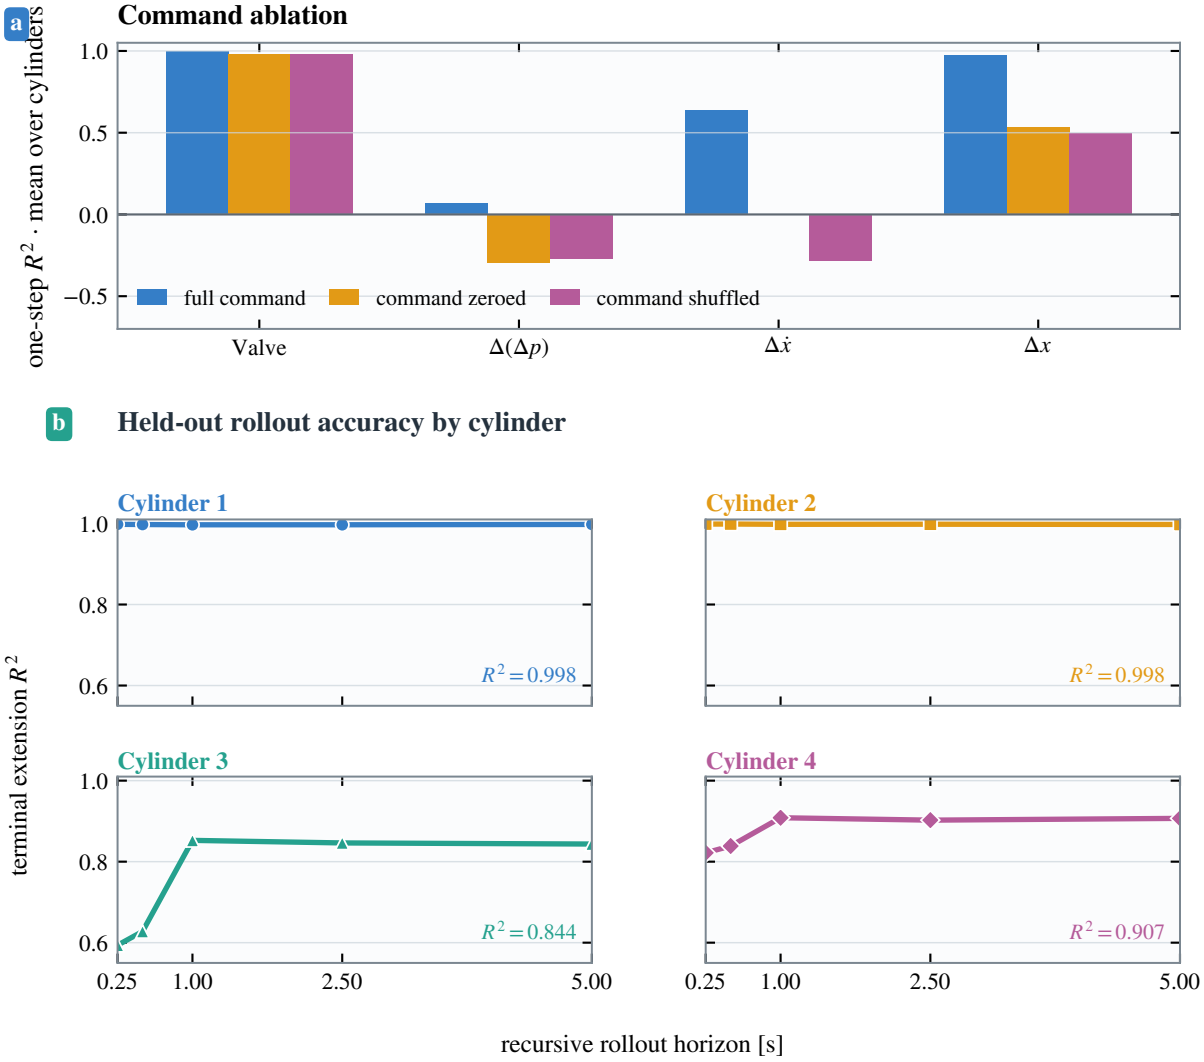

Figure 5.5: Offline validation of the command ROM on held-out runs. (a) One-step R^2 per output group, averaged over the four cylinders, with the full desired command, with the command zeroed, and with the command shuffled across the batch. The collapse of the velocity and extension scores without the true command demonstrates command authority. (b) Terminal extension R^2 of recursive rollouts up to 5 s, shown separately for each cylinder so coincident high-accuracy traces remain visible.

5.4.1 Closed-Loop Calibration Against the Deployment Plant

The Simscape-trained ROM has a systematic extension offset when driving PhysX. The promoted checkpoint adds a bounded, command-gated correction:

$$\hat{x}_{i,k+1} \leftarrow \hat{x}_{i,k+1} + \delta_i(\boldsymbol{h}_k, \boldsymbol{c}_k), \qquad |\delta_i| \le 3\,\mathrm{mm}, \qquad \delta_i\big|_{\boldsymbol{c}_k=\boldsymbol{0}} = 0, \tag{5.9}$$

It is trained on 4 084 IsaacLab transitions from five reachable reference shapes, with Simscape replay. All hydraulic and integration parameters remain frozen. The gate sets the correction to zero for a commanded hold.

The calibrated command ROM serves as the deployment plant in Section 3.4.1. Unlike the ID–FD path, it has seen PhysX data. It is frozen and has no latent input. The earlier

prototype assessed in Sections 5.5 and 5.6 modifies the inverse path. Final AF instead omits the inverse model and corrects the frozen actor's valve command, as described in Section 6.7.

5.5 Offline Arm Adaptation Module Study

This section describes the encoder trained offline against the plant context of Section 4.11, on the frozen collection of Section 4.7. Its purpose is evidentiary: it establishes that plant identity is recoverable from causal observable history at all, which is a precondition for anything built on top of it. The encoder deployed in AF is trained differently (Section 6.7.5), so the results below measure the information in the history rather than describing the deployed component.

The offline AAM estimates $\mathbf{z}_{\mathrm{AAM}} = \mathrm{AAM}(\boldsymbol{H}_t)$ from deployment history. Each row contains five joint positions, five joint velocities, four cylinder pressure differences and four previous valve commands. A four-layer, width-64 causal convolution processes 20 rows (1.0 s, 360 values).

The pressure values are recursive actuator-model states, not measurements: they come from the command ROM in R and the frozen FD in AF. The factor vector, true latent, reference, future rows and hidden state are excluded. Using the previous valve command keeps each row causal. Validation selected the shortest competitive history from 5, 10, 20 and 40 ticks. Latent RMSE is 0.802 on validation and 0.734 on reserved test.

Inference runs at 20 Hz inside the 60 Hz control loop, the prototype ID adapter is bypassed until 20 rows are available, and a causal five-prediction rolling mean adds 250 ms lag. That smoother belongs to this offline encoder and not to AF, whose residual consumes its history window directly, so the latency stated here does not apply to the adaptive result of Chapter 7.

5.5.1 Information Content of the Estimated Latent

Linear decoding from the per-run mean estimate gives $R^2 = 0.956$ for seal condition, 0.448 for compliance, 0.114 for flow, 0.081 for supply pressure, 0.056 for centre-of-mass offset and -0.008 for mass. The latent therefore represents mainly seal condition. It is stable within a run but does not refine over time.

Low-dimensional adapters can lose information when factors act through different channels [Liu 25]. This differs from the rapid-motor- adaptation assumption that only task-relevant factor directions need to be retained [Kumar 21]: here the dominant seal-condition direction acts through pressure, which the controller does not integrate. Estimating named physical quantities would improve interpretation but requires an explicit plant model [Hu 26].

5.6 ID-Adapter Prototype

The earlier prototype applies the correction to the inverse model's output, the direct adaptation point while the inverse model remains in the command path. The closed-loop evidence below shows it has too little authority, and the prototype is retained because that result determines the redesign of Section 6.7.

The frozen ID produces $\boldsymbol{u}_{\text{base}} = \text{ID}(\boldsymbol{s}, \boldsymbol{r})$. The adapter adds a bounded latent-conditioned residual:

$$\boldsymbol{u} = \text{clip}\Big(\boldsymbol{u}_{\text{base}} + g_\psi(\boldsymbol{s}, \boldsymbol{r}, \boldsymbol{u}_{\text{base}}, \mathbf{z}_{\text{AAM}})\Big). \tag{5.10}$$

Training minimises valve-command MSE at the exact ROM tick, with the residual bounded to 0.15, and the ID, FD and policy are frozen. This route uses a separately trained 68-input latent-conditioned policy, so disabling the ID residual restores the baseline actuator path but not the baseline actor. Final AF removes that asymmetry (Section 6.7.3).

Figure 5.6 shows the causal convolutional stack and the bounded adapter that consumes its estimate. It documents the earlier prototype, not the final AF architecture. The final history-residual path is shown separately in Figure 6.6.

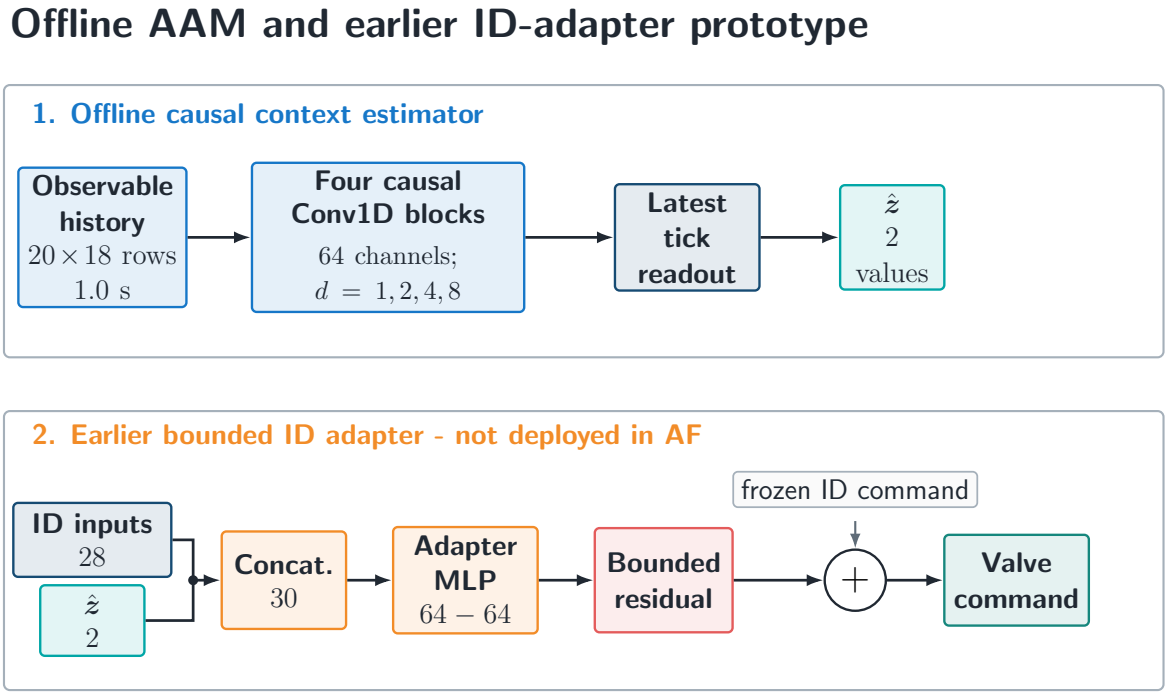

Figure 5.6: Earlier offline AAM and ID-adapter prototype. Four causal convolutional blocks map a 20-step history of 18 observable features to the two-value estimate $\hat{\boldsymbol{z}}$. The bounded adapter combines this estimate with the ID inputs and corrects the frozen ID valve command. This prototype is retained as diagnostic evidence; deployed AF omits the inverse model and trains its history encoder and actor residual jointly under PPO.

5.6.1 Open-Loop Evidence

The residual cap is 0.15, and the frozen ID already explains about 84 % of valve-command variance (Section 5.2.2). The adapter can therefore make only small corrections. Section 5.6.3 tests latent information without this ceiling.

Three matched adapter arms use a zeroed, estimated or privileged latent. Across five seeds, both latent arms improve validation loss over the zeroed arm in 5/5 cases. Mean

paired improvements are 0.0037 for the privileged latent and 0.0118 for the AAM estimate. The effect is consistent but small.

Masking pressure and force does not strengthen the result: the privileged latent helps in only 2/5 seeds and variance increases by about one order of magnitude.

5.6.2 Closed-Loop Evidence and Its Bound

Closed-loop ablations compare the full stack with zeroed, shuffled, frozen and mean latents, plus a sham plant change. None consistently improves lateral error, cut-height error or impacts.

Valve saturation limits the test: cylinder 1 is saturated for about 92 % of the scored window, cylinder 2 for 67–82 %, and cylinder 4 for 50–57 %. Control-effort ratios before and after the plant change are 1.005–1.009. A bounded additive correction has little authority under these conditions [Sofrony 24].

Reducing the lead cap lowers cylinder-1 saturation from 0.91 to 0.14 but triples cut-height error. The task is therefore close to the plant capability limit (Section 3.6.2).

5.6.3 Propagation of the Latent Advantage

The chain gate compares $SD_0(\boldsymbol{s}, \boldsymbol{u})$, $SD_z(\boldsymbol{s}, \boldsymbol{u}, \mathbf{z}_{\mathrm{SD}})$ and $SD_z(\boldsymbol{s}, \boldsymbol{u}, \mathbf{z}_{\mathrm{AAM}})$ under the frozen protocol with paired bootstrap intervals. Neither the true nor estimated latent improves extension or velocity over SD_0 with an interval excluding zero. The estimate is not worse than the true latent, so estimation is not the limiting stage. The measurable advantage remains in pressure, as it was for the separately trained pair of Section 4.11.

The estimate is therefore not the main loss. The missing link is a path from pressure information to the integrated extension command with sufficient actuator headroom.

Every stage of the chain up to the command holds: plant identity is present, a causal estimator recovers it, the estimate is as informative as a privileged one, and it improves a model that consumes it. Delivery fails. The correction is added to a command already resting against its own clip for most of the scored window, so the authority to act on the information does not exist at the point where the information arrives. A limit of that shape is not addressed by a better estimator or a larger latent, nor by reducing the demand, since lowering the lead cap buys headroom at the cost of tripling cut-height error (Section 5.6.2).

It is addressed by removing the accumulator that consumes the headroom, which is what AF does: the inverse model leaves the command path, the policy emits the valve command itself, and the residual acts on that command (Section 6.7.1). The result of that redesign is reported in Section 7.5; Chapter 8 returns to what the sequence implies for adaptation on saturating actuators generally.

5.7 Training and Model-Selection Protocol

Adaptation data are split 82/20/16 into training, validation and reserved test by complete run and factor configuration; windows do not cross run boundaries. The 16 test runs are not used for model or hyperparameter selection. SD_z and SD_0 are selected on validation

rollout, the offline AAM on downstream rollout, and the earlier ID adapter on validation command error. Table 5.3 records the selected checkpoints.

Table 5.3: Provenance of each model in the modelling chain. Frozen components are loaded unchanged and never receive gradients in this chapter.

Model	Status	Selected checkpoint
Control ID	frozen	paired-control, Section 5.2
Strict FD	frozen	strict-excitation, Section 5.3
Command ROM	frozen	PhysX-calibrated, Section 5.4.1
SD_0 (no latent)	scratch	epoch 58
SD_z + encoder	scratch	epoch 48, frozen after the gate
Offline AAM	scratch	epoch 73, 1.0 s causal CNN
ID adapter	scratch	epoch 116

The ID adapter is trained and selected on recorded valve-command error. Each component is evaluated on its stated training contract. The adaptation residual of AF is deliberately absent from Table 5.3, because it has no supervised checkpoint: its encoder and head are optimised jointly under the reinforcement objective, and their provenance is a training run rather than an epoch number.

With the models identified, audited and frozen, the next chapter builds the controllers that consume them. It assembles the reference and non-adaptive stacks from the models above and introduces the final adaptive stack described in Section 5.6.3.

6. Control Framework and Simulation Integration

Chapter 5 left three actuator representations, each with its own contract and offline evidence. This chapter puts a controller in front of them and turns them into three experiments.

Section 6.1 describes the classical controller, because it is where the competence comes from even though it never runs at deployment. Section 6.2 specifies the policy: what it sees, what it emits, what stops it from emitting something dangerous, and what it is paid for. Section 6.3 explains how the first becomes the second, and reports a result about that transfer which changed the training protocol of this thesis. The individual stacks appear only after those three are fixed.

6.1 The DLS Teacher and Its Two Roles

6.1.1 Scope and Role of the Controller

No damped-least-squares layer runs underneath the policy at runtime and nothing is added to the policy's output; if the policy emits zero, the arm does not move. The DLS controller is nevertheless the reason the policies work at all, and it appears twice, both times at training or collection time and never at deployment (Figure 6.1). It is the *privileged teacher* whose behaviour the policies are distilled from (Section 6.3), and it is the closed-loop controller that drove the identification runs of Section 4.7. Calling it a nominal layer would misdescribe the deployed system and would make every reported result a property of a controller pair rather than of the policy.

6.1.2 Privileged Information

The teacher is *privileged* rather than better tuned: it reads two quantities that no deployable controller on this machine could obtain, namely the exact terrain height directly under the cutter, queried from the simulator rather than inferred from a forward-looking scan,

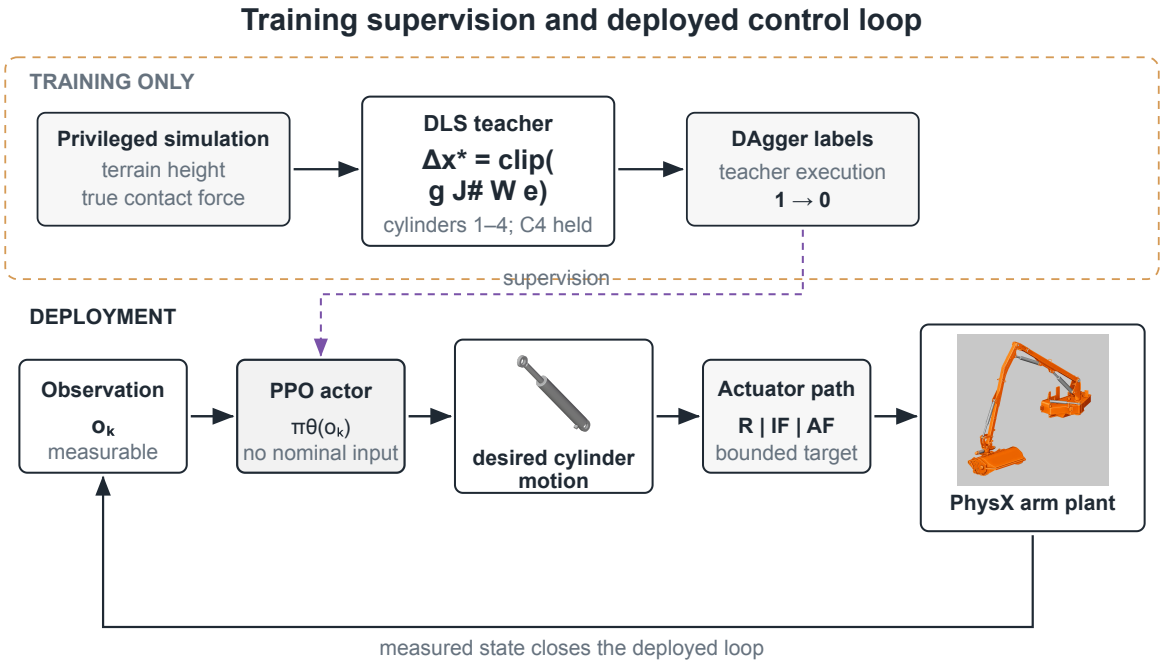

Figure 6.1: Training supervision and deployment. Simulator-only terrain and contact signals drive the DLS teacher, which supplies DAgger labels to the actor. At deployment only the lower feedback loop remains: measured observations pass through the PPO actor, the shared cylinder interface and the selected R, IF or AF actuator path before PhysX updates the arm. No nominal command is added to the policy output.

and the true cutter–ground contact force. Together these turn the hardest part of the task, knowing where the ground is underneath a 1.2 m flail head that is occluding it, into a lookup.

The teacher uses that privilege in a specific way. Rather than commanding a height, it runs a contact regulator: it maintains a vertical offset relative to the ground and adjusts that offset to hold the measured contact force near 250 N within a 100 N deadband, pressing down below the band, backing off proportionally above it, searching downward when contact is lost and retreating immediately on a force spike. The accumulated offset is clamped to $[-0.30, -0.05]$ m so the regulator can neither drive the head into the ground nor lift it clear.

The teacher's advantage is therefore not kinematic. Its inverse is the same DLS solution available to anything; what it has is contact sensing, and that is the quantity the student must infer from a noisy forward scan and its own contact history.

6.1.3 From Task Error to Cylinder Command

The rest of the teacher is conventional and is stated for completeness. A goal point is maintained in the truck frame: laterally it tracks the verge centre line, vertically it follows the measured ground height plus the contact regulator's offset. The task error is the difference between that goal and the measured cutter position, with the lateral component

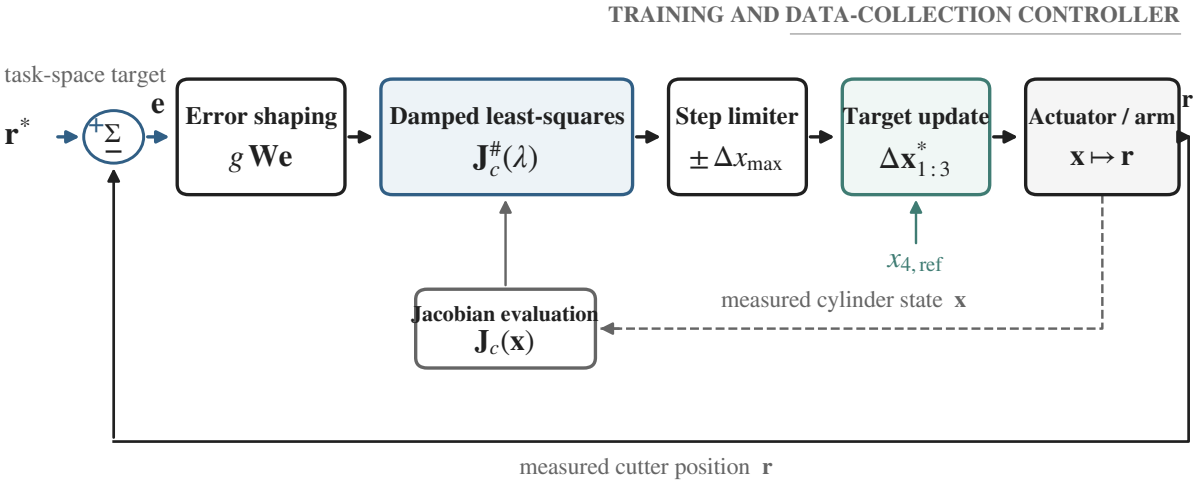

Figure 6.2: Feedback realization of the damped least-squares teacher of (6.1). The weighted task error is scaled by g, mapped to the first three cylinder coordinates by the measured damped inverse $\boldsymbol{J}_c^{\#}(\lambda)$, and clipped to $\pm\Delta x_{\max}$ per tick. Cylinder 4 is held at its head-down reference. Cutter position closes the task-space loop and the measured cylinder state updates $\boldsymbol{J}_c$ at every step. The loop is used only for teaching and data collection, never at deployment.

weighted by a factor of three, because on this machine lateral error is what the acceptance gate measures and vertical error is largely handled by contact.

That error is mapped to cylinder space through the damped least-squares inverse of the measured cutter Jacobian $\boldsymbol{J}_c$ of Section 3.2.2,

$$\begin{aligned} \boldsymbol{J}_c^{\#}(\lambda) &= \boldsymbol{J}_c^{\mathsf{T}} \left(\boldsymbol{J}_c \boldsymbol{J}_c^{\mathsf{T}} + \lambda^2 \boldsymbol{I} \right)^{-1}, \\ \Delta \boldsymbol{x}^*_{1:3} &= \operatorname{clip}\Big(g\, \boldsymbol{J}_c^{\#}(\lambda)\, \boldsymbol{W} \boldsymbol{e},\ \pm \Delta x_{\max} \Big), \quad g = 0.30, \quad \Delta x_{\max} = 3\,\mathrm{mm}. \end{aligned} \tag{6.1}$$

with $\boldsymbol{W}$ the lateral weighting and the per-step clamp bounding how far any cylinder may be asked to move in one tick. Only cylinders 1–3 are solved. Cylinder 4 is held at the reference extension that keeps the head down, for the reason given in Section 3.1. The gain g is small on purpose: this is a stiff, heavy mechanism in contact, and a resolved-rate controller that closes quickly on the task error simply excites it.

6.1.4 Execution Through the Student Interface

One implementation detail determines whether distillation converges to something usable. The teacher's output $\Delta \boldsymbol{x}^*$ is not written to PhysX directly: it is converted into the same normalised four-value action the student emits and pushed through the student's entire action chain into the same plant, so every label the student imitates is by construction an action it can produce. A teacher writing targets directly to the plant would generate labels outside the student's reachable set, giving the imitation error a floor no training could remove.

6.1.5 Teacher Performance

Quoting the teacher's performance is necessary because distillation cannot exceed it and the gates of Section 3.9 were anchored to it. Table 6.1 reports the privileged teacher

Table 6.1: The privileged DLS teacher, measured on the stage-6 verge task through the command-ROM plant (128 environments, 480 steps, all eleven path families, seed 17). It is competent, not perfect, and several acceptance thresholds sit deliberately just outside what it achieves.

Quantity	Teacher	Gate
Mean $\lvert e_{\text{cut}} \rvert$	0.0234 m	≤ 0.040 m
p95 $\lvert e_{\text{cut}} \rvert$	0.1524 m	—
Mean $\lvert e_{\text{lat}} \rvert$	0.1221 m	≤ 0.300 m
Ground-contact fraction	0.7807	≥ 0.700
p01 head alignment	0.9638	≥ 0.900
Minimum head alignment	0.8317	—
p95 tool acceleration	64.9 m/s^2	≤ 120 m/s^2

measured on the stage-6 verge task through the ROM plant, over 128 environments and 480 control steps.

The teacher is competent but visibly imperfect: it loses ground contact about a fifth of the time and its cut-height tail is six times its mean. This is why Section 3.9 re-derived several thresholds, since a gate set above these numbers could not be passed by the ground-truth controller itself and would carry no information about the policy.

6.2 Policy Formulation

6.2.1 Policy Definition

The deployed controller is a memoryless Gaussian multilayer perceptron. It maps a 66-value observation to four normalised actions, one desired motion command per hydraulic cylinder, and that is the entire control law. Actor and critic are separate networks with hidden widths 256–128–64 and ELU activations, over empirically normalised observations. This description is of the base policy, and it is shared by all three stacks: the actor width is the same 66 in each, and AF deploys this network frozen and unmodified rather than a widened variant of it. Two things differ for AF and are treated where they arise. Its four actions are interpreted as valve commands rather than as desired cylinder increments (Section 6.7.2), and the deployed controller is this policy plus an additive residual whose own encoder reads a history window carried alongside the observation, leaving the base actor's input unchanged (Section 6.7.3).

The base policy is not itself a residual and has no nominal command to correct. The AF residual is added around this recoverable frozen base. Temporal inference is assigned to a separate encoder (Section 6.7.5), while the truck's forward motion remains scripted.

6.2.2 The Observation Vector

Figure 6.3 lays out the 66 channels by group. The grouping tells the story of what this controller has to reason about: nine terrain-height samples ahead of the cutter, three task-error channels, sixteen channels of cylinder state and remaining travel, sixteen describing the actuator interface, ten of pose and context, and twelve of model error.

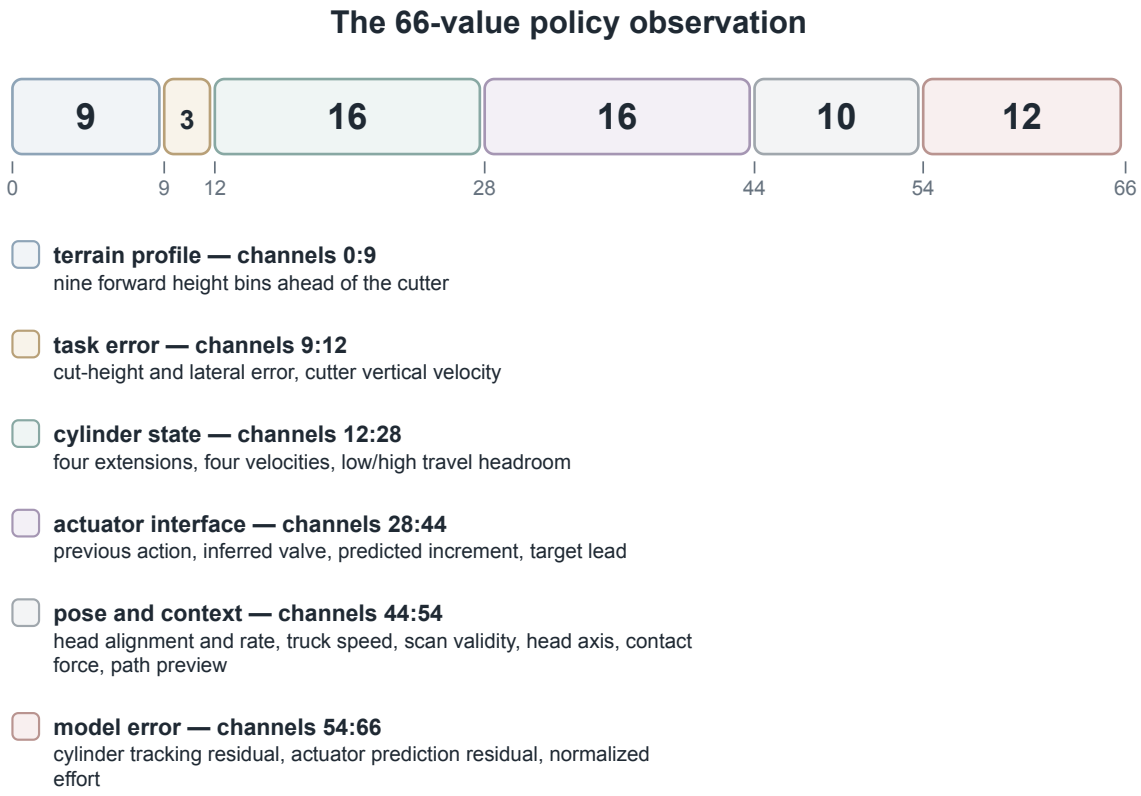

Figure 6.3: The 66-value policy observation, drawn to scale by group. The channels combine measurements, derived quantities and actuator-model estimates. No trajectory-family label, physical factor vector or latent estimate enters the actor. AF keeps this 66-channel input unchanged and carries its causal history separately.

The last group of twelve channels describes the actuator rather than the task, carrying its standing error as the cylinder tracking residual, the actuator prediction residual and the normalised cylinder effort. They exist as a direct consequence of the policy being memoryless. Forming an error requires differencing position across two steps and comparing against a prediction issued on the previous step, and a feedforward network can construct neither from a single observation, so without these channels the model error the policy is asked to compensate is absent from the input rather than merely hard to see. This was measured: across ten training blocks the p01 head alignment sat at 0.79–0.80 and did not respond to changes in the feedforward source or the levelling gain.

This group is also what makes the matched pair comparable. Under the ROM the tracking residual is near zero, because the ROM reaches what it is told, while under a learned ID–FD path it is the model's standing error. The same observation contract therefore gives each policy an honest view of its own actuator's shortcomings, which is what Section 6.9.2 requires.

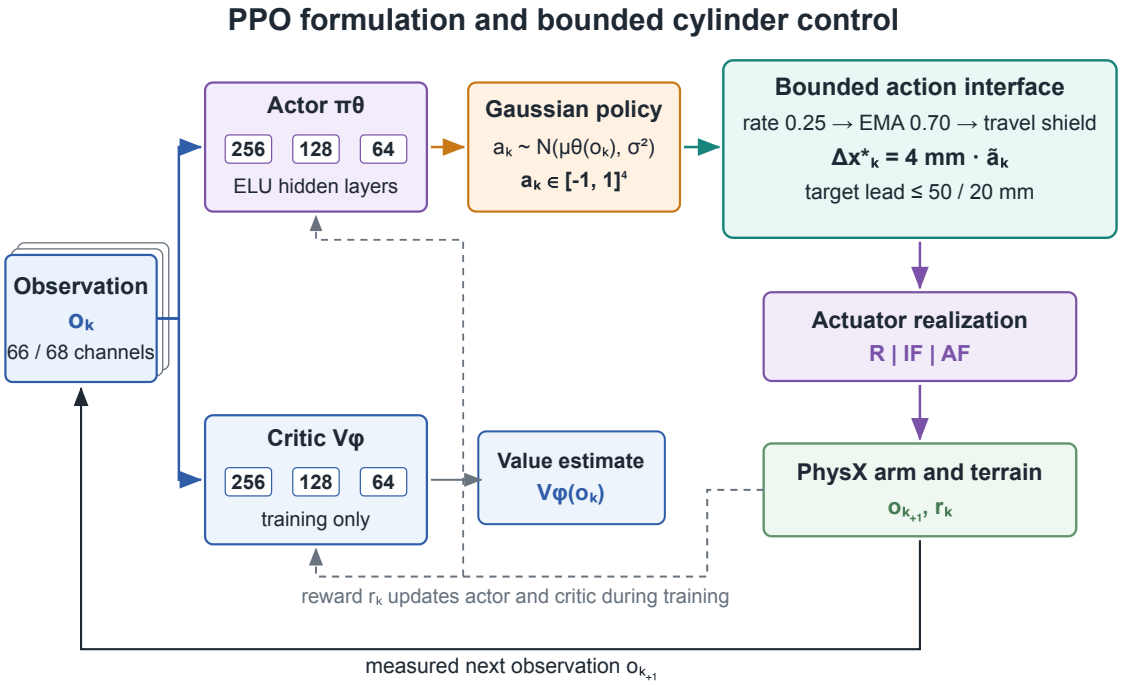

Figure 6.4: PPO actor–critic formulation and deployed control path. The actor maps the observation to a four-dimensional Gaussian action. Rate limiting, exponential filtering, the travel shield, command scaling and the target-lead cap bound that action before the selected R, IF or AF actuator path drives the PhysX arm. The critic and reward update are training-only; measured state closes the deployed loop.

6.2.3 The Action and Its Bounds

The policy emits $\boldsymbol{a} \in \mathbb{R}^4$, sampled from its Gaussian at training time and taken at the mean for evaluation. Between that sample and the PhysX cylinder target sit five bounding stages. Figure 6.4 places them in the complete PPO actor–critic loop. The critic is used only during training; the actor and bounded action interface form the deployed controller. Each bounding stage can only *reduce* authority and cannot introduce motion the policy did not request.

The rate limit bounds $|\Delta \boldsymbol{a}|$ at 0.25 per control step and is held *constant* across every curriculum stage rather than used as a curriculum knob, by the reachability argument above: the teacher's own command stream has a structural per-step bound near 0.20, so a student limit below that makes part of the label set unreachable. The exponential filter ($\beta = 0.70$ at the final stage) smooths the command. The limit shield scales the command down, and only in the direction that consumes headroom, over the last 20 mm of cylinder travel, replacing the teacher's hard clamp into the learned range. The command scale converts the filtered action into 4 mm of cylinder travel per unit action, and the target integrator accumulates the increments in reference space, permitted to lead the measured extension by at most 18 mm.

This chain describes the cylinder-increment interpretation of the action, which R and IF use. Under AF the same bounded action is read as a valve command and scaled to span exactly the admissible valve range, so no further clamp binds at nominal. The bounding

stages that act in cylinder space therefore do not have the same meaning on that stack, which is one reason its results are read against its own frozen base (Section 6.9.2).

That last cap is an anti-windup device and it is load-bearing. Relaxing it from 18 mm to 50 mm halved the lateral error on the two families demanding a sustained standing offset (`lane_change_hold` from 0.4833 to 0.2310, `boom_sweep` from 0.3745 to 0.1602), showing that the clamp was bounding the standing effort those paths need, while the same relaxation on cylinder 4 drove head alignment from 0.9507 to 0.4636. The final task therefore resolves the trade per axis, and this is the configuration every reported R result was produced under: 50 mm on cylinders 1 to 3 and 20 mm on cylinder 4.

The chain replaces the residual band, since exploration is bounded by how much motion a single action can command rather than by proximity to a baseline, and it fixes the plant's capability ceiling: four millimetres of cylinder travel per step is the entire authority budget, which produces the 0.240 m/s lateral rate limit of Section 3.6.2 from which the reference paths are derived.

6.2.4 The Reward Function

The reward of the final R task has ten bounded terms, listed in Table 6.2. Three are positive and describe the task. Seven are penalties and describe how the machine must behave while doing it. All are bounded, and all retain a live gradient beyond the acceptance boundary, so that a policy already inside a gate still has a direction to improve in.

It replaces an inherited formulation carrying 53 weighted reward fields, at which size the optimum is not something a designer can reason about: three separate terms turned out to be working against the task without being noticed for thirteen training campaigns. Three rules govern the replacement, each adopted after a measured failure, and Appendix D records the measurements behind them.

No Gaussian Shaping

The earlier reward priced tracking as $w\exp[-(e/s)^2]$, which is effectively flat beyond about $2s$ and therefore supplies almost no gradient where the failing path families live. Every shaping term is instead linear or Huber in the error. Writing $x = |e|/s$, the two shapes used are

$$B(e;s,c) = \min(x,\, c), \qquad H(e;s,c) = \min\Big(c,\, \begin{cases} \tfrac{1}{2}x^2, & x \le 1 \\ x - \tfrac{1}{2}, & x > 1 \end{cases}\Big). \tag{6.2}$$

B has the same gradient at 0.5 m of error as at 0.05 m, so a policy far from the reference is still told which way to move; H buys quadratic precision near zero where cut height needs it, without a vanishing far field.

Bounded Terms

A term that can grow without limit can dominate the sum and destabilize the critic. Each contribution is clipped, which gives the reward a published budget: $6 + 12 + 4 = 22$ for the task terms against -24 for the safety terms. A non-finite value is replaced by -25, just below that worst case, so a numerical fault can never outscore legitimate behaviour.

Table 6.2: The ten reward terms of the final R task, with weights and shaping scales. B and H are the bounded shapes of Equation (6.2). The lateral and head terms each combine a coarse and a fine scale, so far-field correction and near-field precision are priced separately.

Term	Weight	Shape and purpose
Cut-height tracking	+6	H, scale 0.04 m, cap 2.5; the primary mowing-quality term
Lateral tracking	+12	B, coarse 1.50 m and fine 0.60 m mixed 0.6/0.4; the primary line-following term
Ground contact	+4	pays for the cutter being in contact
Head alignment	−12	B on the deficit below 0.98, fine scale 0.02 and coarse scale 1.00 mixed evenly
Contact loss	−3	charges directly for loss of contact
Action rate	−1	B on the mean absolute command change, scale 0.25; penalizes chatter
Tool acceleration	−1	B, scale 300 m/s^2; mechanical-wear proxy on the cutter
Contact force	−1	B on force above 30 kN, scale 150 kN
Cylinder headroom	−2	B on encroachment inside a 0.02 m margin, scale 0.10 m
Hard impact	−4	indicator on the impact event

Fixed Sign Convention

The three task terms are non-negative and the seven safety terms non-positive. This is what makes worst-case (CVaR) weighting safe to add later: scaling a non-positive term by a factor ≥ 1 can only lower the return, whereas up-weighting a *positive* reward for the worst environments pays the policy to be bad, as measured in the earlier stack.

Three terms of the earlier formulation were removed on measured evidence rather than preference, and Appendix D records the figures. The *action-magnitude* penalty on mean(a^2) charged for precisely the sustained offset that `lane_change_hold`, `offset_return` and `boom_sweep` structurally require, so the policy learned to return to neutral, which is the one thing "hold an offset" forbids. The *head barrier* was replaced because the head group dwarfed lateral tracking, and the *clearance barrier* was disabled because measured contact, cylinder headroom and cut-height error proved more reliable signals of the same condition.

One apparent inconsistency is deliberate. The tool-acceleration term saturates at 300 m/s^2 while the corresponding gate sits at 120 m/s^2. The reward scale is a saturation scale for a bounded penalty rather than a target, so the gate remains the binding criterion and the reward's role is only to keep a gradient against gross excursions. The action-rate term is held mild for a related reason: cylinder 4 owns head pitch, so a rate tax falls hardest on the fast corrections that keep the head down as boom attitude shifts.

Every term is logged individually and by name during training, because the whole difficulty is that total reward can rise while the system gets worse. The reward is therefore *not* the result criterion: promotion uses the eight physical gates of Section 3.9.

6.2.5 Optimisation

Policies are optimised with clipped-surrogate PPO over generalized advantage estimation, in the form of Section 2.7. The maintained configuration uses clip parameter $\epsilon = 0.2$, discount $\gamma = 0.99$, GAE parameter $\lambda = 0.95$, five learning epochs over four minibatches per update, an adaptive learning-rate schedule targeting a KL divergence of 0.02 from a base rate of 3×10^{-4}, gradient-norm clipping at 1.0, and a clipped value loss with coefficient 1.0. The terrain task collects 32 steps per environment per update. At 512 environments that is 16 384 transitions per update, chosen because the inherited 128-step horizon exceeded two minutes per update on the available GPU without adding useful batch diversity.

The final fine-tuning blocks depart from that configuration because they run under the reference anchor of Section 6.4 rather than an adaptive learning-rate schedule: clip $\epsilon = 0.1$, discount $\gamma = 0.97$, three learning epochs, fixed learning rates of 10^{-5} for the actor and 10^{-4} for the critic, entropy coefficient 0.001, an exploration standard deviation fixed at 0.15 and excluded from the gradient, and frozen observation normalisers.

Three of those departures are structural rather than tuning. The discount is reduced because the episode is long in simulation steps but locally controlled, so an action should be credited for correcting the terrain state it faces rather than for a return hundreds of steps later that the critic cannot resolve. The actor learning rate sits an order of magnitude below the critic's so that the actor is deliberately the slow variable. And the adaptive schedule is dropped because the anchored formulation already constrains the update through an explicit divergence bound; two competing step-size controllers would make neither attributable.

The exploration settings are also specific to this architecture. A zero action here means the arm does not move at all rather than that it tracks nominally, so the initial standard deviation was raised from 0.05 to 0.30 to let the policy discover the command directions before refining them. On the terrain task, where distillation supplies a competent starting point, it is reduced again and *held fixed* rather than initialized, for the reason taken up in Section 6.4. Exact values for any reported run are read from that run's frozen manifest (Table D.2).

6.3 Training Protocol and the Distillation Finding

6.3.1 Distillation by DAgger

Policies are initialized by imitation, not by reward. The privileged teacher is rolled out in the student's own environment and the student is fitted to the teacher's raw action by regression, in the DAgger scheme of Ross et al. [Ross 11]: the teacher's probability of *executing* the action, rather than merely labelling it, is annealed from one to zero over the first 70 % of iterations, so later training collects states the student itself visits. Labels are accumulated in a replay buffer of 262 144 transitions and refitted over three epochs per iteration in minibatches of 8 192.

Figure 6.7 shows a complete distillation run for the ROM stack, taken from the stored record. The behaviour-cloning loss falls by more than an order of magnitude while the teacher is still driving, reaches its minimum around the handover, and then *rises* once the student is alone. That rise is the distribution shift DAgger exists to expose rather than a

failure: the student is now visiting states the teacher never chose, and the teacher's advice on those states is new information. The closed-loop task metrics remain near the teacher's preflight values without matching them.

6.3.2 Objective Mismatch in Unconstrained Fine-Tuning

Unconstrained PPO fine-tuning from distilled checkpoints increased total reward while degrading the primary tracking objective, in the worst case by a factor of 12.9. The critic remained healthy, with explained variance in the 0.85–0.95 range. This motivates the anchored formulation of Section 6.4, which is used for the final policies.

The reward includes ground contact and cut height as well as lateral tracking. It can therefore increase while a gate-relevant tracking measure degrades.

This is a claim about two functionals over the same policy space rather than about the optimiser. Training maximizes the expected return $J(\vartheta)$, while acceptance is decided by the vector of physical gates $\boldsymbol{G}(\vartheta) \preceq \boldsymbol{g}^\star$ of Section 3.9. Writing $\mathcal{F} = \{\vartheta : \boldsymbol{G}(\vartheta) \preceq \boldsymbol{g}^\star\}$ for the feasible set, the measured content of the finding is

$$\arg\max_{\vartheta} J(\vartheta) \notin \mathcal{F}, \qquad \text{while} \qquad \vartheta_{\text{distilled}} \in \mathcal{F}. \tag{6.3}$$

Distillation lands inside $\mathcal{F}$ because the teacher was constructed to satisfy the gates, not to maximize J. Unconstrained ascent on J from that point therefore leaves $\mathcal{F}$ by construction. This is the optimiser working correctly on the wrong functional, and it is why the remedy is a constraint on the ascent rather than a repair of the reward.

Two consequences follow, both of them protocol. Unconstrained fine-tuning is inadmissible, and both the pre- and post-fine-tune checkpoints are reported in every case. And reward decomposition, critic explained variance, policy entropy and the gate metrics themselves are logged *during* training rather than only at the end, since a training curve reporting only the scalar reward would have shown a success here.

The general form applies beyond this machine: when a competent teacher exists and the acceptance criterion is not the reward, reward-based fine-tuning is not automatically an improvement step. Distillation transfers the intent, PPO then optimises the proxy, and the two point in measurably different directions. That divergence grows with distance travelled from the policy where proxy and intent were measured to agree, so the remedy is to bound the distance explicitly rather than to abandon reward optimisation.

6.4 Reference-Anchored PPO

Equation (6.3) says that reward ascent leaves the feasible set. It does not say that reward ascent is useless, only that it is unbounded. The construction of this section keeps PPO and constrains how far it may travel from a policy already known to satisfy the gates. Writing π_{ref} for a frozen reference policy and J for the expected return, each training block solves

$$\max_{\vartheta} J(\vartheta) \quad \text{s.t.} \quad \mathbb{E}_{\boldsymbol{o}\sim d_{\pi_\vartheta}}\Big[\mathrm{KL}\Big(\pi_{\text{ref}}(\cdot \mid \boldsymbol{o}) \,\|\, \pi_\vartheta(\cdot \mid \boldsymbol{o})\Big)\Big] \le \delta_{\text{ref}}, \tag{6.4}$$

with the expectation taken over the states the policy itself visits.

6.4.1 Critic Warm-Up

After distillation the actor already behaves sensibly, but the critic has never been fitted to the long terrain return, so the first PPO updates would move a useful actor using a value estimate that is not yet meaningful. The actor is frozen for 150 iterations while only the value loss is optimised, under the asserted invariant $\vartheta_{\text{after}} = \vartheta_{\text{before}}$. The observation normalisers are frozen at the same point, since otherwise the same raw observation is normalised differently from one iteration to the next and the actor's behaviour changes while its weights do not.

6.4.2 The Reference Anchor

Let π_{ref} be a frozen copy of the accepted warm actor. It is queried only on states the live policy visits and never supplies an executed action, so PPO remains the sole controller of the plant. The penalty is the *forward* divergence, which for diagonal Gaussians is available in closed form,

$$D_{\text{ref}} = \mathrm{KL}\big(\pi_{\text{ref}} \,\|\, \pi_\vartheta\big) = \sum_i \left[\log \frac{\sigma_{\vartheta,i}}{\sigma_{\text{ref},i}} + \frac{\sigma_{\text{ref},i}^2 + (\mu_{\text{ref},i} - \mu_{\vartheta,i})^2}{2\sigma_{\vartheta,i}^2} - \frac{1}{2}\right], \tag{6.5}$$

at the cost of one additional forward pass and no sampling. Because both scales are fixed at $\sigma = 0.15$, this collapses to a pure quadratic on the drift of the commanded action,

$$D_{\text{ref}} = \sum_i \frac{(\mu_{\text{ref},i} - \mu_{\vartheta,i})^2}{2(0.15)^2}, \tag{6.6}$$

which is the second reason for fixing σ: with a free standard deviation, entropy collapse shrinks D_{ref} while the commanded action moves, so the constraint would report compliance during the drift it exists to prevent.

The direction is load-bearing. Forward KL is mass-covering: it charges π_ϑ wherever π_{ref} has density and π_ϑ does not, and therefore penalizes the *abandonment* of reference behaviour. The reverse direction is mode-seeking and would permit the policy to collapse onto a subset of reference behaviour at no cost, which is the observed failure, in which one competence was dropped to buy another. It follows that the term is a behaviour-cloning anchor evaluated on the policy's own state distribution, rather than the reverse-KL regularizer of soft reinforcement learning. Its lineage is demonstration-regularized policy gradient [Rajeswaran 18], not the Boltzmann-optimum family.

The actor objective is

$$\mathcal{L}(\vartheta) = \mathcal{L}_{\text{PPO}} + \mathcal{L}_V - \beta\,\mathcal{H}(\pi_\vartheta) + \lambda_{\text{ref}}\,\mathbb{E}_k\Big[\mathrm{KL}(\pi_{\text{ref}}\|\pi_\vartheta)\Big], \qquad \lambda_{\text{ref}} = 0.10. \tag{6.7}$$

6.4.3 Trust Regions and Rollback

A penalty is not a constraint: at a fixed coefficient, a sufficiently large advantage always buys its way past. Feasibility is therefore also verified on the update taken. Two bounds are imposed,

$$D_{\text{update}} = \mathrm{KL}(\pi_{\text{old}}\|\pi_\vartheta) \le 0.01, \qquad D_{\text{ref}} \le 0.05, \tag{6.8}$$

and they do different work. D_{update} is the classical step-size control that keeps the surrogate a valid local model of J. D_{ref} is the new object: because the anchor is *fixed* within a block rather than moving with the policy, it bounds cumulative drift. A sequence of updates each individually satisfying D_{update} can still walk arbitrarily far from feasibility, and that is the trajectory the earlier campaigns took, monotonically, with the reward rising the whole way. Only a constraint referred to a fixed feasible point can observe it.

Enforcement is in two stages. During minibatch optimisation the epoch is early-stopped as soon as D_{update} exceeds its bound, and after the update both divergences are recomputed over the full rollout. If either bound fails, or any quantity is non-finite, the actor and its optimiser state are restored *atomically*. On a divergence rejection the critic fit is retained when finite, since the bounds constrain the policy rather than the value function; on a non-finite update both networks are restored. In the final block, five of one hundred logged updates were rejected, all on the reference bound and none on the update bound, so the drift caught was of the kind ordinary PPO clipping does not see.

6.4.4 The Block Ratchet

The reference is not fixed forever. Within a block it is frozen. When a checkpoint is scored and passes the gates, it is promoted to reference,

$$\vartheta_{\text{ref}}^{(m+1)} = \vartheta^{\star}, \qquad \vartheta^{\star} \in \mathcal{F}, \quad \vartheta^{\star} \in B_{\delta_{\text{ref}}}\left(\vartheta_{\text{ref}}^{(m)}\right), \tag{6.9}$$

and the actor's optimiser moments are reset, because momentum accumulated relative to the old anchor is not meaningful relative to the new one. The critic's are retained: it is fitting the same return, and its calibration is expensive to rebuild.

This is the structural point of the construction. The policy is at all times within a bounded divergence of a checkpoint that has been verified against $\boldsymbol{G}$, and the anchor advances only through verified checkpoints. The acceptance criterion therefore enters the training loop instead of appearing only at the end. The optimiser still ascends the proxy. It simply can never travel far from the last point at which proxy and intent were measured to agree. In this construction distillation is no longer the endpoint: it is $\vartheta_{\text{ref}}^{(0)}$.

Algorithm 6.1 summarises the procedure in five stages. The first two produce the initial verified anchor. Stages 3 and 4 repeat for each PPO block, and Stage 5 returns the last promoted checkpoint.

6.4.5 Theoretical Status

The components are standard and citable. The bound on D_{update} is conservative policy iteration and its trust-region descendants [Kakade 02, Schulman 15]; the atomic rollback is that literature's backtracking line search reduced to a single trial point; the penalty in Equation (6.7) is a Lagrangian relaxation of the constrained problem (6.4) at a fixed multiplier, with the rollback supplying the feasibility enforcement a fixed multiplier cannot.

What does not follow is a certificate. Bounding the divergence from a verified policy does bound gate degradation in form, since the chain rule of KL over trajectories, Pinsker's inequality and the boundedness of every gate metric give $|\mathbb{E}_{\pi_\vartheta}[g] - \mathbb{E}_{\pi_{\text{ref}}}[g]| \leq 2G_{\max}\sqrt{H\delta_{\text{ref}}/2}$ for a metric g with $\|g\|_\infty \leq G_{\max}$ over a horizon H. At the production horizon of 2808

Algorithm 6.1: Five-stage reference-anchored PPO procedure. The teacher is used only in Stage 1. All PPO rollouts are controlled by π_ϑ alone, and the frozen reference moves only when a checkpoint passes every gate.

input : teacher π_T; gate test $\boldsymbol{G}(\vartheta) \preceq \boldsymbol{g}^\star$; M PPO blocks; anchored loss and divergence bounds from Equations (6.7) and (6.8)
output : last promoted actor $\vartheta^\star$ satisfying the gate test

`// Stage 1 - Distil a feasible actor`
$\mathcal{D} \leftarrow$ DAgger states labelled by teacher mean action $\boldsymbol{y}^T$
$\vartheta \leftarrow \arg\min_\vartheta \mathbb{E}_{(\boldsymbol{o},\boldsymbol{y}^T)\sim\mathcal{D}}\big[\|\mu_\vartheta(\boldsymbol{o}) - \boldsymbol{y}^T\|_2^2\big]$
remove π_T from the controller `// all later rollouts are teacher-free`
`// Stage 2 - Warm the critic and set the first anchor`
freeze ϑ and the observation normalisers
for $j \leftarrow 1$ **to** 150 **do**
  roll out π_ϑ and set $\psi \leftarrow \psi - \eta_{\text{critic}} \nabla_\psi \mathcal{L}_V(\psi)$
end
verify $\vartheta_{\text{after}} = \vartheta_{\text{before}}$
$\vartheta_{\text{ref}} \leftarrow \vartheta$; $\vartheta^\star \leftarrow \vartheta$ `// initial verified anchor`
for *block* $m \leftarrow 0$ **to** $M - 1$ **do**
  `// Stage 3 - Train inside the frozen reference bound`
  freeze ϑ_{ref} for the whole block
  for *each PPO iteration in block* m **do**
    collect an on-policy rollout with π_ϑ as sole controller
    compute and normalise GAE; save snapshot $\mathcal{S}$; set $\vartheta_{\text{old}} \leftarrow \vartheta$
    for 3 *epochs over the rollout minibatches* **do**
      if $\mathrm{KL}(\pi_{\text{old}} \| \pi_\vartheta) > \delta_{\text{upd}}$ **then** stop the update epochs
      update actor and critic using Equation (6.7)
    end
    compute $(D_{\text{upd}}, D_{\text{ref}})$ on the full rollout `// full rollout`
    if D_{upd} *or* D_{ref} *is non-finite* **then**
      restore the complete snapshot $\mathcal{S}$
    else if $D_{\text{upd}} > \delta_{\text{upd}}$ ***or*** $D_{\text{ref}} > \delta_{\text{ref}}$ **then**
      restore the actor and its optimiser moments from $\mathcal{S}$; retain the new critic fit
  end
  `// Stage 4 - Evaluate the block and promote a passing checkpoint`
  score every saved checkpoint on at least three held-out terrain meshes
  aggregate each gate to the worst mesh and, for tracking gates, the worst path family
  $\mathcal{C} \leftarrow \{\vartheta_c : \boldsymbol{G}(\vartheta_c) \preceq \boldsymbol{g}^\star\}$
  if $\mathcal{C} \neq \varnothing$ **then**
    choose $\vartheta^\star \in \mathcal{C}$ using the fixed checkpoint selection rule
    load $\vartheta \leftarrow \vartheta^\star$ and set $\vartheta_{\text{ref}} \leftarrow \vartheta^\star$
    reset the actor optimiser moments; retain the critic state
  else
    leave ϑ_{ref} unchanged; do not promote the block
  end
end
`// Stage 5 - Return the deployable policy`
return $\vartheta^\star = \vartheta_{\text{ref}}$

steps and $\delta_{\text{ref}} = 0.05$ the right-hand side exceeds eight, which for a metric bounded in $[0, 1]$ is vacuous, and there is a further measure mismatch because the constraint is taken under the live policy's state distribution while the chain rule requires the reference's. The construction is therefore a principled restriction of reward ascent to a neighbourhood of a verified policy, not a guarantee: the gates certify each promotion.

6.5 R: Reference Actuator Baseline

R is the reference condition of this thesis. The policy emits desired cylinder motion, the frozen command ROM converts it to a cylinder position target, and PhysX closes the loop. There is no online adaptation anywhere in the stack. The truck is scripted over a 26 m strip at one constant speed drawn per episode from 2 to 7 km/h, the eleven path families of Section 3.6 are distributed across environments within a single run, and the terrain mesh is drawn per environment so that no policy can memorize one ground surface.

Its status is that of a reference condition rather than a competitor. The command ROM is a Simscape-trained model with a bounded closed-loop calibration against IsaacLab and PhysX, and ROM-plus-PhysX *is* the deployment plant, so R's actuator coincides with the plant. The consequence for how its tracking advantage may be read is stated once, in Section 6.9.2.

The policy is produced by the three-layer construction of this chapter: distillation from the privileged DLS teacher to a feasible student, critic warm-up with the actor frozen, and reference-anchored PPO in blocks under Equations (6.7) and (6.8). The final block produced checkpoints 280 to 329. The last iterate was not taken, because PPO can raise mean return while damaging a single path family. Candidates were scored on held-out meshes against the eight gates, aggregated to the worst mesh and, for tracking-critical gates, the worst path family. Chapter 7 reports the final common-protocol evaluation.

6.6 IF: Non-Adaptive Learned Baseline

IF subclasses the R task and changes exactly one thing. The policy still emits desired cylinder motion, but that command is now realised by the frozen inverse model of Section 5.2, which converts it to a valve command, followed by the frozen strict forward model of Section 5.3, which converts the valve command to a predicted extension increment. PhysX receives the same bounded cylinder target as before. Task, terrain families, meshes, observation contract, action chain, reward, gates and the PPO update are inherited unchanged, so the actuator representation is the only free variable between the two stacks.

The stack is the honest control for R, whose actuator coincides with the deployment plant, and it is the stack in which the inverse model is carried in the command path and is adequate there.

Robust is used here in the sense of Section 2.6: the stack was trained across the randomised plant family and applies one fixed response to whichever member of that family it meets, performing no online estimation and conditioning on no plant identity. It does not mean the policy is blind to plant variation, since the twelve model-error channels carry a standing residual that is itself plant-dependent. What the stack lacks is not information about the consequences of plant variation but any representation of *which* plant produced them.

Two properties of the composed path shape its behaviour, and both were established offline. The valve command is accumulated rather than issued absolutely and is clipped to the per-axis limits $[0.95, 0.95, 0.80, 0.80]$, so the actuator path carries its own saturation boundary independently of the action chain. Only the predicted extension increment is integrated; the velocity output is retained as a diagnostic (Section 5.3).

The stored run manifest shows that the actor learning rate, entropy coefficient, seed, terrain seed, environment count, curriculum stage, frozen normaliser, and warm-up match R field for field (Table D.2). This makes the pair a controlled comparison. Chapter 7 reports the common-protocol results.

Figure 6.5 summarises the two non-adaptive actuator paths used as baselines. Both use the same policy interface and PhysX task; only the actuator representation changes.

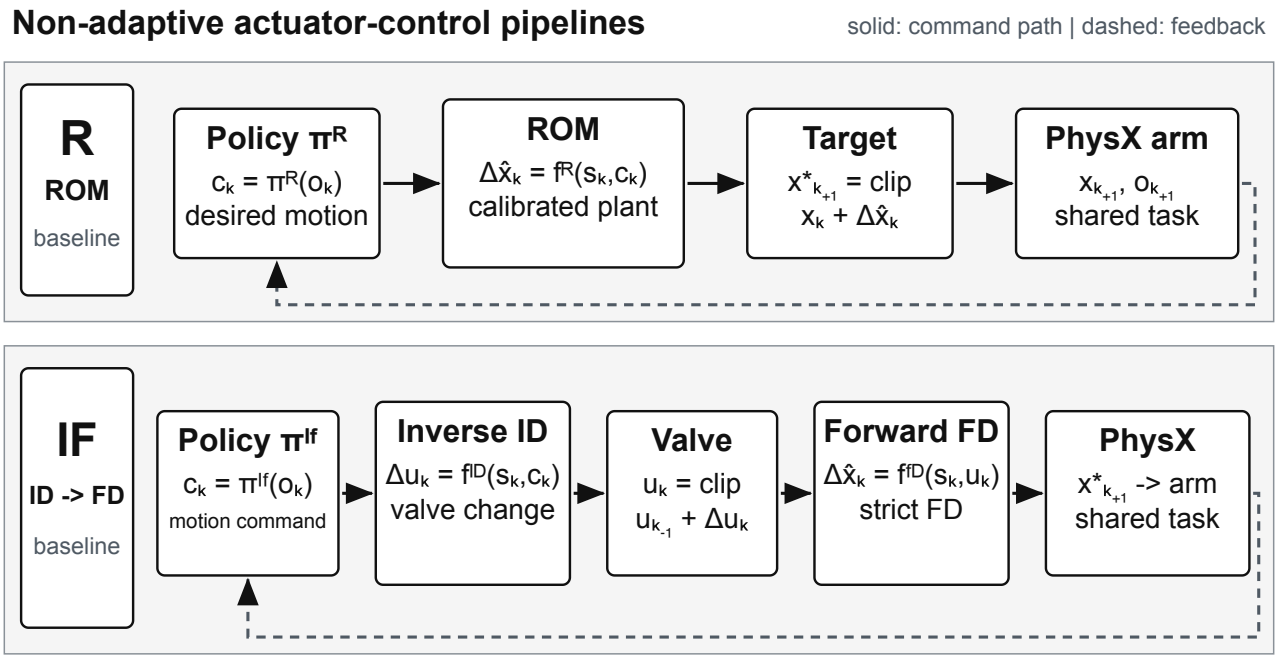

Figure 6.5: The non-adaptive R and IF control pipelines. R sends the policy's desired cylinder motion through the calibrated command ROM. IF composes the independently identified inverse and strict forward models before applying the same bounded cylinder target in PhysX. Both paths use the same observation, action limits and task.

6.7 AF: The Adaptive Stack

AF places online adaptation on the forward model alone. The policy emits the valve command directly, the frozen strict forward model of Section 5.3 converts that command to a predicted extension increment, and PhysX receives the same bounded cylinder target as the other three stacks:

$$\pi_{\text{frozen}} + \text{history residual} \longrightarrow \boldsymbol{u} \longrightarrow f_{\text{FD}} \longrightarrow \text{PhysX}. \tag{6.10}$$

The inverse model is absent from the command path. This is the one structural difference from IF, and it is deliberate: the name of the stack records the two models that remain,

AAM and FD. Task, terrain families, meshes, observation contract, reward, gates and the PPO update are inherited unchanged, so the actuator representation and the adaptation path are the only free variables against the baselines.

That the inverse model is genuinely out of the loop is a verified property rather than a claim about intent. Under the direct-valve configuration the adaptation-site instrumentation reports zero valve elements attributable to the inverse model and zero adapter applications at that site across a complete run, where all three counters would be non-zero if the inverse model were in the path (Appendix C.3). The actuator key recorded inside the checkpoints still reads `id_aam_fd`, retained so that previously trained checkpoints remain loadable; it is a legacy identifier, not a description of the runtime path.

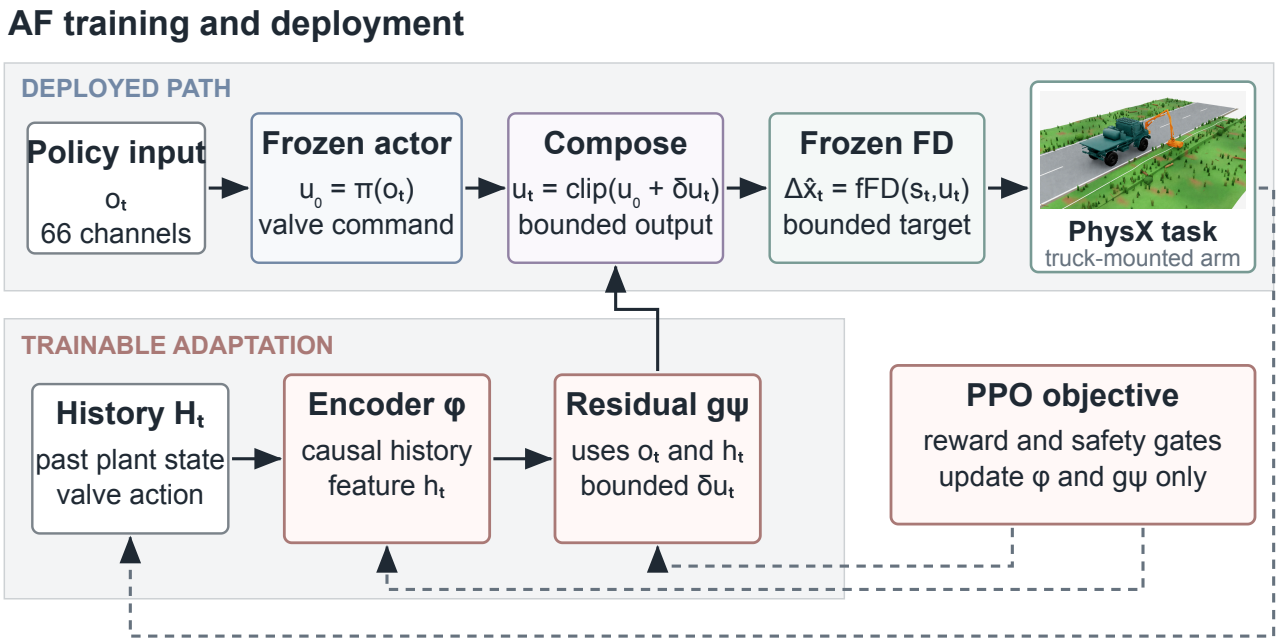

Figure 6.6: Training and deployment path of AF. The frozen base actor emits the valve command $\boldsymbol{u}_0$. A temporal encoder and residual head use causal state–action history to add the bounded correction of Equation (6.11). Only these two modules receive the PPO task gradient; the base actor and strict forward model remain frozen. The composed valve command drives the forward model and the shared PhysX verge-following task. The actor input remains 66 channels; the history window is carried separately and is not part of that count. Dashed paths denote PPO updates and rollout feedback. No privileged factor vector or inverse model enters the AF command path.

6.7.1 Omission of the Inverse Model

The inverse model's place in two of the three stacks follows from what each stack is for. The command ROM together with PhysX *is* the deployment plant (Section 5.4.1), so R uses the plant itself and fixes the reference condition. The inverse and forward models are identified independently of it, from the high-fidelity Simscape data of Chapter 4, and neither has seen the deployment plant during training, so IF prices the cost of substituting an independently identified pair for the plant itself. That question requires the inverse model: a policy commanding desired cylinder motion needs a map from that intent to

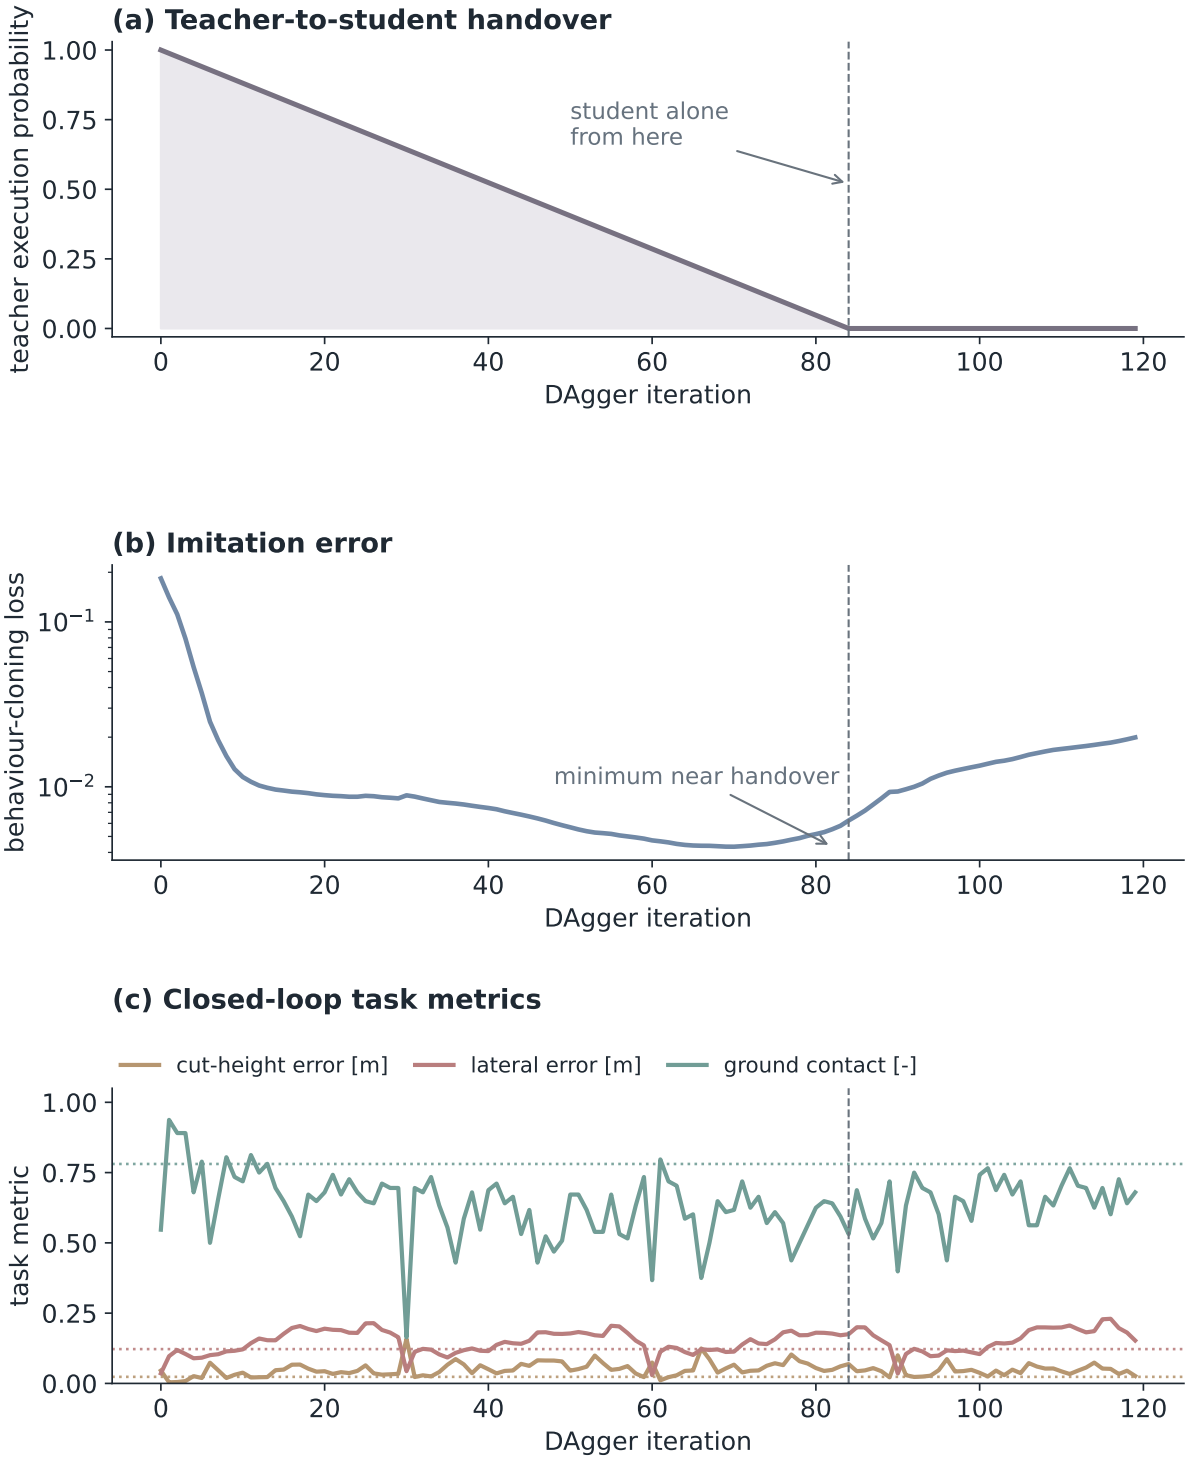

Figure 6.7: A measured distillation run for the R stack. (a) The teacher's execution probability falls to zero over the first 70 % of iterations. (b) Behaviour-cloning loss falls by over an order of magnitude under teacher execution, then rises as the student visits new states. (c) Cut-height error, lateral error and ground-contact fraction; dotted lines mark the teacher's preflight values from Table 6.1.

a valve command, and supplying that map from identification data alone is exactly the substitution being priced. The R–IF gap is interpretable only because the composition is complete.

The adaptive stack asks a different question. The inverse model remains appropriate for IF and its one-step transfer is established in Section 5.2.2. In the adaptive command path, however, an additional correction must retain authority after the command is formed.

The same transfer study recorded the failure mode that later disqualified it for adaptation. The model predicts a valve-command *increment*, and the deployed command is the accumulated, clipped quantity $\boldsymbol{u}_k = \mathrm{clip}(\boldsymbol{u}_{k-1} + \Delta\boldsymbol{u}_k, -\boldsymbol{u}_{\max}, \boldsymbol{u}_{\max})$ with per-axis limits [0.95, 0.95, 0.80, 0.80], so under sustained demand the accumulation walks the command onto its own clip and holds it there. Fed its own predictions, the composite left the reference by more than a metre in *every one* of the 118 runs, with its pre-clip valve command at the limit on 89.1 % of elements while the ROM valve on the same runs never reached its bound. Closed-loop instrumentation over the 33-cell panel measures the same thing in the terms that matter for adaptation: the command sat at its pre-clip limit on 79.9 % of valve elements, 34.5 % of the corrections the adapter computed were clipped away after being computed, and on 37.9 % of elements the correction could not move the command at all.

The quantity being adapted and the quantity being saturated are therefore the same, and the consequence is not a small effect but an uninterpretable one: a null measured through such a path cannot distinguish a task that is not adaptation-limited from credit that was clipped before it could reach the metric. The disqualification is specific and retracts nothing claimed for IF, where the command is consumed as issued and saturation is a limit on authority rather than a defect of the map. It binds only where a second signal must be added to that command and survive the same clip.

The final stack therefore emits the valve command directly. The residual acts on the same bounded command as the policy, rather than on an increment that is subsequently accumulated and clipped.

6.7.2 Sufficiency of the Forward Model

Removing one of the two learned actuator models raises the question of whether the remaining one carries the path. Sufficiency here is a specific claim: the forward model must convert a commanded valve position into a one-step extension increment with genuine valve sensitivity, because that increment is the only quantity the controller integrates.

The evidence is the strict retraining of Section 5.3.1. Before it, zeroing the valve input at a 0.5 s horizon changed the rollout-error ratio to only 0.94, meaning the network had extrapolated smooth closed-loop kinematics rather than learned the valve's effect. After strict retraining on the open-loop plateau data the valve-zero error ratio rises to 1.81, extension rollout RMSE falls from 75.2 mm to 12.4 mm at 0.5 s and from 157.9 mm to 27.3 mm at 1.0 s, and per-axis one-step extension R^2 improves from [−5.7, 0.60] to [0.76, 0.98].

Sufficiency is established on the channel that is used and is not claimed beyond it. Three qualifications carry forward. Only the extension increment is integrated, the velocity output remaining weak on two axes and retained as a diagnostic. Zero-order hold remains

competitive below 0.25 s, so the forward model earns its place at the horizons the controller actually spans rather than uniformly. And the valve authority the model represents is strongly operating-point dependent, varying by a factor of 15–75 across pressure quartiles and capable of changing sign, which is why a constant valve Jacobian was rejected.

6.7.3 The Residual Controller

Adaptation enters as an additive, bounded residual on the valve command:

$$\boldsymbol{u} = \mathrm{clip}\Big(\boldsymbol{u}_0 + \alpha \tanh\Big(g_\psi(\boldsymbol{o},\, \phi(\boldsymbol{H}_t))\Big),\, -1,\, 1\Big), \qquad \boldsymbol{u}_0 = \pi_{\mathrm{frozen}}(\boldsymbol{o}), \tag{6.11}$$

where ϕ is a temporal-convolution encoder over a window of (state, action) pairs and g_ψ is a small multilayer perceptron producing a per-cylinder correction. Only ϕ and g_ψ are trained: 53 836 parameters against a 59 588-parameter frozen base.

Three construction choices make the comparison in Chapter 7 a single-variable one, and each was verified rather than assumed. The final layer of g_ψ is zero-initialised and $\tanh(0) = 0$, so at construction the residual is exactly zero and the wrapped policy reproduces the frozen base bit for bit, at a measured maximum absolute action difference of 0.0. Every arm of an ablation therefore begins as the identical policy, and any end-state difference is attributable to the residual having been learned rather than to the task having been relearned. The correction is bounded by $\alpha \tanh(\cdot)$ regardless of what the head produces: with head weights drawn from $\mathcal{N}(0,\, 50)$ the output saturates at exactly α and remains finite, so a bad update degrades accuracy instead of emitting a non-finite action, which matters because a non-finite action written into the history buffer would poison the entire window. And the residual scale can be set to zero at evaluation time without touching any other part of the stack, in which case the stack reduces to the frozen base exactly.

This last property improves on the earlier latent-conditioned design, whose actor was itself a 68-input network, so that disabling the residual restored the baseline actuator path but not the baseline actor and the ablation was never quite single-variable. Under AF the base actor is the distilled policy unmodified, the residual carries the whole of the context dependence, and disabling it recovers the baseline in full.

The history window is carried in the observation vector rather than in a buffer held inside the module, because PPO calls the actor many times per update epoch and a module-held buffer would be advanced by the optimisation loop and desynchronise from the simulation. The causal convention is fixed throughout: the estimate at step t uses $\boldsymbol{x}_{t-k:t-1}$ and $\boldsymbol{a}_{t-k:t-1}$ and does not include step t.

6.7.4 Freezing the Base Policy

Freezing the base is structural rather than a tuning preference, and the alternative was closed by measurement before the residual route was adopted. Under a trajectory-regularised fine-tune with a reference-KL anchor the policy is stable but cannot acquire a new input dependence, because the anchor's reference is context-blind by construction, so any context-dependent behaviour increases KL to that reference and is penalised. A gradient multiplier of 1000 applied to the context columns alone moved action sensitivity only from 0.001 % to 0.004 % of action scale over 400 iterations, and masking the anchor

at $\boldsymbol{z} = \mathbf{0}$ reached 0.005 %. The anchor is not the blocker, however: without it PPO walks away from the distilled behaviour, trading line tracking for ground contact because that raises its objective.

A residual head has no anchor. It is new capacity whose only function is to depend on the latent, so nothing opposes that dependence. Measured on the same stack, the trained residual's action responds to its history at 3.35 % of action scale, roughly 840 times the anchored route.

6.7.5 Training of the Adaptation Encoder

The encoder ϕ and the head g_ψ are trained together under the reinforcement objective, not fitted offline to an auxiliary target and handed to a controller afterwards. This is the concurrent scheme of Ji et al. [Ji 22], and its consequence is that the latent is control-relevant by construction rather than by hope: the only gradient reaching ϕ is the one that improves the task.

An offline dynamics objective does not by itself guarantee a control-relevant history representation. The adaptation module is therefore assessed behaviourally in Chapter 7, through action sensitivity to history and the closed-loop comparison with the frozen base.

6.7.6 Consistency With the Plant Dynamics

Two consistency requirements follow from the system model of Chapter 3, and both are met by construction.

The first is causality. Each history row carries observable quantities only, namely joint positions, joint velocities, cylinder pressure differences and the previous valve command. The privileged factor vector, the true latent, the reference, future rows and any hidden simulator state are excluded, and the estimate at step t is formed without step t, so nothing enters the deployed path that would be unavailable on the physical machine.

The second is a qualification rather than a guarantee. Under AF the pressure differences supplied to the encoder are recursive states of the frozen forward model, not measurements of the plant, so the encoder observes the plant through the forward model's representation of it and its estimate is consistent with the true dynamics only to the extent that the forward model is. Section 6.7.2 bounds that extent. This limit on the interpretation of the latent is carried into Chapter 7 rather than resolved here.

The stack does not claim a physically named estimate. The residual is trained to improve the task, and the quantity ϕ extracts is whatever serves that objective; it is not constrained to correspond to seal condition, compliance, supply pressure or any other named factor. Recovering named physical quantities would require an explicit plant model and a different estimator, noted as future work in Chapter 8.

6.7.7 Scope of the Claims

The stack is complete and its interfaces are verified. The residual authority α is a genuine tuned hyperparameter and is reported as one rather than presented as a natural constant. The deployed value is $\alpha = 0.5$. At $\alpha = 0.15$ the residual is not merely less effective

but actively worse than the base it corrects, by 28.3 % on cut-height error paired across 528 environments, with the interval on the paired difference excluding zero. A bounded correction needs enough authority to matter, and below that threshold it contributes noise rather than adaptation; no value of α should therefore be read as incidental. Whether the adaptation path contributes anything to closed-loop task performance is a separate question, settled by the paired comparison of Section 7.5 and not asserted here.

6.8 IsaacLab Integration and Timing

Four rates run concurrently and Table 3.4 lists them: PhysX integrates at 120 Hz, the policy acts at 60 Hz under a decimation of two, the learned actuator models advance on a 20 Hz grid of 0.05 s model ticks, and the AAM updates on that same grid. No single one of these is "the control frequency", and reporting a scalar rate for this system would misdescribe every one of the three stacks.

The relationship between the 60 Hz and 20 Hz layers is where integration errors hide, so it is fixed by construction rather than by interpolation. A learned actuator model is evaluated on every third control step and its predicted increment is held across the two intervening steps. The action presented to the model is the action of the tick on which it fires, not an average, and the target integrator accumulates in reference space, because integrating against the measured extension turns a static hold into a slow ratchet (Chapter 3).

Order within a tick is fixed and asserted at run time. The history encoder consumes the window ending at the previous tick and produces its estimate before the actuator path runs, so no quantity computed at the current tick can influence its own estimate. History buffers are per-environment and cleared on episode boundaries, since the vectorised environments reset asynchronously and a shared buffer would leak one episode into the next. Both properties are checked rather than assumed: a runtime check replays a trace and asserts bit-exact agreement between the online estimate and an offline causal recomputation. Checkpoint loading then compares the stored normalisation statistics, feature contract and history length of every learned component against what the environment expects, and aborts on a mismatch.

6.9 Constraints and Fair-Comparison Rules

6.9.1 Limits and Safeguards

Authority is bounded at four independent places, and the redundancy is deliberate because each bound protects against a different failure. The action chain of Section 6.2 limits the rate of change of the command, fades authority as cylinder travel headroom is consumed, and caps how far the integrated target may lead the measured extension. The actuator path clips the valve command to its per-axis limits. The environment terminates an episode on the safety conditions of Section 3.9 rather than letting it accumulate reward in a state an operator would stop. And every learned component falls back to a defined behaviour on a non-finite output.

Preflight runs before any training or evaluation and refuses to start on a violated precondition. The check that does the most work asserts that the commanded reference stays inside the measured actuator authority ceiling of Section 3.6.2: the cylinders move

the cutter laterally at most 0.240 m/s, so a path family demanding more is infeasible by construction and no reward, critic or optimiser change could recover it. Terrain is derived from that ceiling, and the preflight assertion keeps the two from drifting apart.

6.9.2 Matching Rules and Comparison Controls

This thesis does not present a three-way ranking of the stacks. Two distinct comparisons are made, each with its own matching rule and each carrying one claim; conflating them would produce a table whose rows are not commensurable.

The Matched Pair: R Against IF

The representation claim rests on R and IF, which agree on every recorded training field and differ only in the actuator path. Held identical across the pair are the verge task and its eleven path families, the terrain generator and the held-out mesh set, the reward function of Table 6.2, the action chain and its authority budget, the eight acceptance gates and their aggregation to the worst mesh and worst path family, episode length, actor and critic capacity, and the three-layer training construction of distillation, critic warm-up and reference-anchored PPO. That pair is a controlled single-variable comparison and is the only place an across-stack number is read as such.

One departure is declared. Because the command ROM plus PhysX *is* the deployment plant (Section 5.4.1), R's tracking advantage over IF is not evidence that a reduced-order model approximates the actuator better than an identified pair: one of the two defines the target.

Non-Comparability of AF and IF

AF does not share the action contract of the other two, since its policy emits the valve command itself where the other two emit a bounded desired cylinder increment that a downstream path realises. The quantity the policy chooses is not the same quantity in the two cases, so the action chain that makes the other comparison fair does not extend to this one. This is a declared property of the implementation, and it is why no row of any table ranks AF against IF on task error. The adaptive claim was never a claim about beating the non-adaptive stack.

The Paired Comparison: AF Against Its Own Frozen Base

The adaptation claim is made within the adaptive stack, against the frozen base policy it was built from, and the matching is exact. The arms share the frozen base policy, the residual architecture, the trainable capacity of 53 836 parameters, the training budget and the reward. Because the head of g_ψ is zero-initialised and $\tanh(0) = 0$, every arm begins as bit-for-bit the same policy, verified at a maximum absolute action difference of 0.0. The single manipulated variable is whether the residual may see the (state, action) history it is conditioned on. A capacity-matched control is trained at every seed, so the comparison is made within a seed rather than across one at all five (Section 7.5.3).

The arms are therefore paired environment by environment, so scenario variation is differenced out rather than averaged over, and any separation between the control and the adaptive arm is attributable to online sensing rather than to added parameters, a different

optimiser trajectory or a different checkpoint. Every arm is scored at the peak of its own training-reward curve by one parser, which uses no held-out information. The result is reported on this basis in Section 7.5.

Fine-Tuning Budgets Must Be Matched or Reported

Comparing a distilled baseline against a separately optimised adaptive policy and attributing the difference to adaptation is not permitted, for the reason established in Section 6.3.2: reward ascent moves a policy off the acceptance criterion, so an unmatched budget is itself a treatment. The paired construction above satisfies this rule by design. The rule is stated because it constrains any future stack added to this comparison, not because the present one is at risk of violating it.

7. Experimental Evaluation and Discussion

This chapter evaluates the closed-loop control stacks introduced in Chapter 6. It addresses the final two research questions:

RQ5 How does the actuator representation affect tracking, command behaviour, latency, and robustness under plant variation?

RQ6 Can recent observations support online adaptation when plant variation affects closed-loop performance?

7.1 Experimental Protocol

7.1.1 Test Conditions and Comparisons

All closed-loop results use held-out terrain meshes, eleven trajectory families, a truck speed of 4 km/h, a 26 m scoring window, and sixteen reserved plant configurations. The matched comparison for RQ5 is between R and IF. These stacks use the same training configuration and differ in their actuator representation. RQ6 is evaluated within AF by comparing a frozen base policy, a capacity-matched residual without history, and a history-conditioned residual.

The results are paired by environment or by mesh and trajectory family. Error metrics are reported separately for lateral tracking and cut height because the two errors describe different aspects of the task. Confidence intervals are obtained by paired bootstrap. Checkpoints are selected using the validation or training criterion specified for each experiment and are then evaluated once on the held-out set.

7.1.2 Comparative Scope and Scoring

R and IF form the cross-stack comparison. AF has a different action contract and is therefore not ranked against that pair. Its results are interpreted only through within-stack comparisons.

Table 7.1: Closed-loop performance over 33 paired held-out conditions. The matched stacks differ only in the actuator representation.

	R	IF
lateral error, mean (m)	0.1026	**0.0703**
lateral error, median (m)	0.0392	**0.0281**
lateral error, p95 (m)	0.4664	**0.2665**
cut-height error, mean (m)	0.0084	**0.0063**
cut-height error, median (m)	**0.0013**	0.0031
ground contact	**0.9623**	0.9529
hard impacts	0.0947	**0.0905**
tool acceleration p95 (m/s^2)	39.90	**39.02**

7.1.3 Scoring Horizon

All reported closed-loop results use the full 26 m trajectory rather than a shorter selection window.

7.2 Closed-Loop Task Performance

The task is scored on two independent error channels, drawn side by side in Figure 7.1: lateral error is deviation from the planned cutting line in the horizontal plane, and cut-height error is deviation from the intended vertical cut. They are kept separate throughout because a controller can improve one while worsening the other. Ground contact, impacts and tool acceleration are reported alongside them as checks on motion quality.

Table 7.1 compares the matched pair over 33 held-out (mesh, family) conditions. IF reduces mean lateral error from 0.1026 m to 0.0703 m and mean cut-height error from 8.4 mm to 6.3 mm. Median errors are low for both stacks. The mean difference is driven mainly by difficult trajectory families rather than by routine tracking. The p95 values show the same pattern, indicating that the improvement also reduces larger tracking deviations.

Figure 7.2 shows that neither stack is best on every trajectory. IF is better on six of the eleven families, while R is better on five. The largest difference occurs on the double S-curve, where IF has substantially lower error. Thus, the inverse and forward model improves average tracking, but the advantage is trajectory dependent. Both stacks follow the simpler trajectories closely. The difficult families provide the clearer test of the actuator representation.

7.3 Actuator Representation and Dynamic Response

7.3.1 Plant Delay and Step Response

Open-loop excitation measured a command-to-pressure delay of approximately 117 ms. Cylinder extension and joint motion began at the same time, which confirms that the dominant delay is hydraulic rather than kinematic. The similar values across excitation amplitudes indicate that this delay is a stable property of the simulated actuation chain.

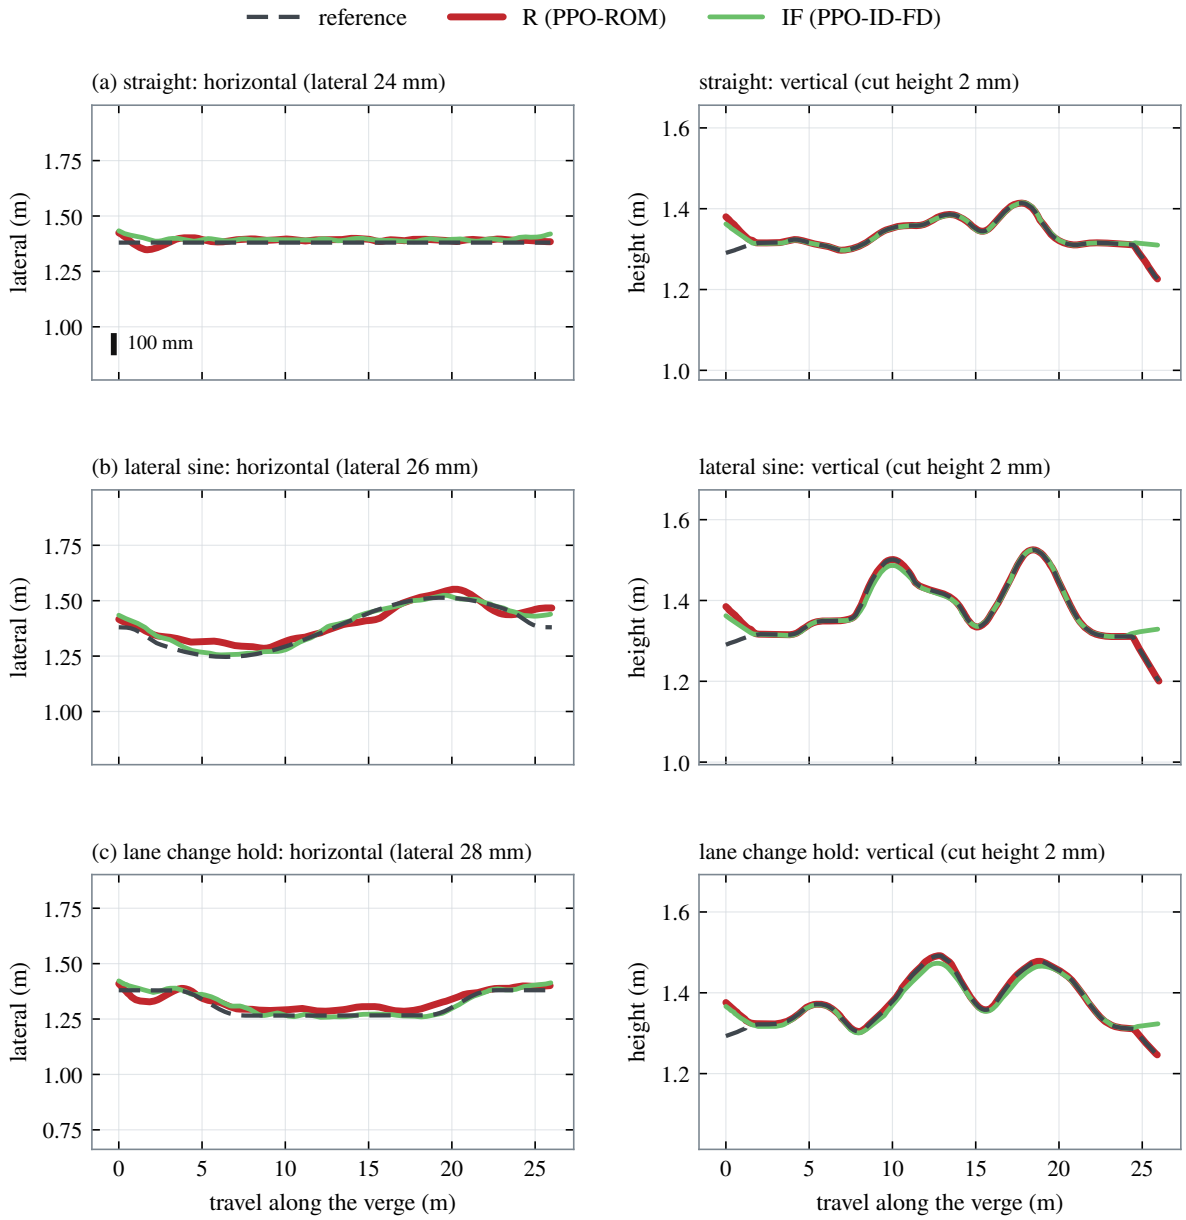

Figure 7.1: Cutter path against the reference on held-out mesh 1337, for three structurally different trajectory families: a flat line, a smooth periodic sweep and a step-and-hold. The left column is the horizontal plane and the right column the vertical plane over the same drive. The vertical reference is not recorded directly and is recovered as $z_{\text{ref}} = z_{\text{cutter}} - e_{\text{cut}}$, so both columns compare reference against actual. In the vertical panels the reference is largely hidden beneath both traces, which is what a millimetre-scale cut-height error looks like drawn at path scale.

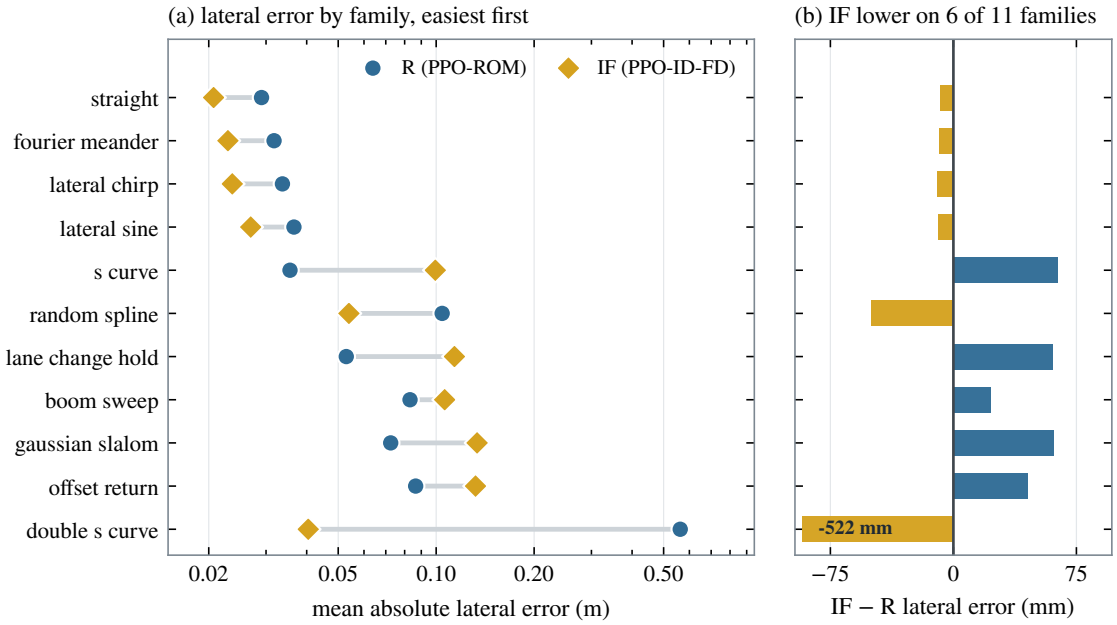

Figure 7.2: Lateral error by trajectory family for the matched pair. The aggregate difference is concentrated in the more difficult families.

Table 7.2: Transport latency along the actuation chain, in milliseconds. Values are medians over open-loop excitation runs at 60 Hz.

	excitation amplitude			
Link	0.05	0.15	0.35	0.70
command → differential pressure	117	117	117	150
command → cylinder velocity	183	167	133	117
command → cylinder extension	117	117	100	100
command → joint position	117	117	100	100

Table 7.3: Response to a 100 mm commanded step. Values are median times over sixteen environments.

Stack	detected response	rise	settled	overshoot
R	333	617	2633	57 %
IF	350	717	2183	73 %

Times are in milliseconds.

Table 7.4: Control effort on a held-out mesh. Lateral error is in millimetres; commands are in normalised units.

Stack	RMS command	rate RMS	reversals/s	lateral error
R	0.586	7.55	40.0	120.1
IF	0.827	10.68	**10.6**	**26.8**

The commanded-step experiment measures the closed-loop response to a 100 mm change in the reference line. Rise time is measured from 10 to 90 % of the step. Settling time is the first point at which the error remains within 20 mm for 0.5 s.

IF settles about 0.45 s earlier, whereas R has lower overshoot. Both responses are much slower than the hydraulic delay. The command rate limit alone requires at least 417 ms to complete this step, and both stacks require more than 2 s to settle. Closed-loop policy dynamics and command limits therefore dominate the response time. The actuator path changes the shape of the response, but it does not determine the main settling time.

This distinction is important for control design. Reducing hydraulic delay alone would not remove the observed settling time. Improvements to the command rate, policy response, and reference handling are also required.

7.3.2 Control Effort

R uses a smaller command magnitude but reverses direction almost four times as often as IF. The reversal count identifies rapid changes in command direction that are not evident from the RMS magnitude alone. IF uses more actuation and achieves lower tracking error with smoother command direction. The actuator representation therefore affects both tracking and command behaviour.

The effort values should be read together with the tracking error. Figure 7.3 draws that pairing: a lower RMS command is not preferable when it is accompanied by frequent reversals and larger path error.

7.4 Generalisation and Robustness

The matched stacks generalise to the held-out meshes, trajectory families, and reserved plant configurations used in this chapter. The randomisation sweep showed little change in tracking over the tested mechanical variation. Its hydraulic interpretation is limited because the inverse and forward plant does not apply the hydraulic factors, while the reduced-order model constrains them to its identification range.

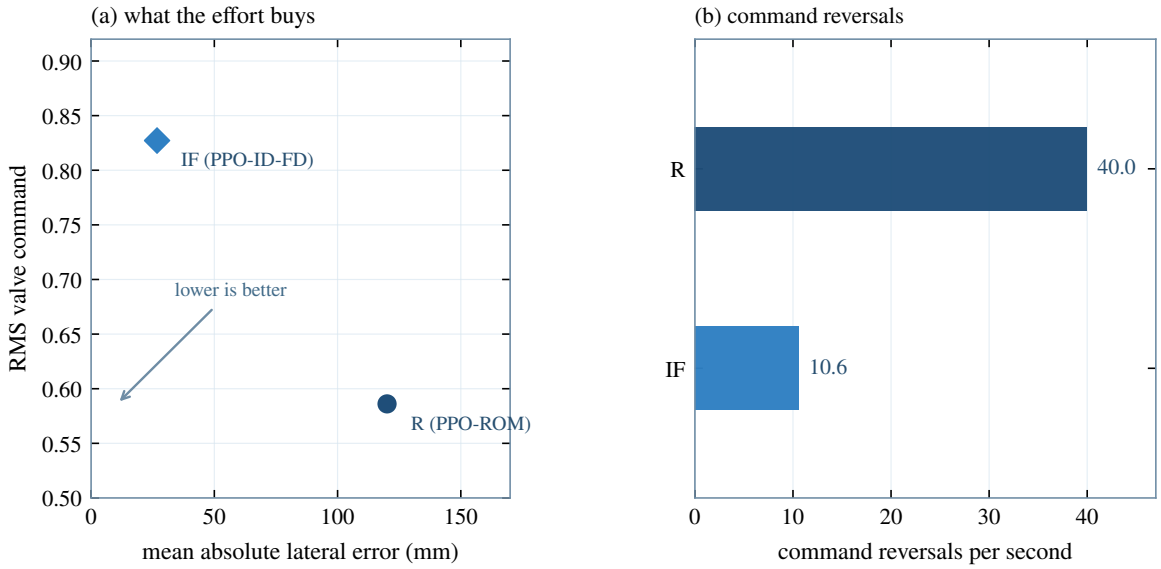

Figure 7.3: (a) Control effort against the tracking error it buys for the matched pair; the lower left is better. (b) Command reversals per second, which the root-mean-square command does not reveal.

The held-out evaluation therefore establishes transfer within the tested simulation population. It does not establish performance on arbitrary hydraulic conditions or physical hardware.

7.4.1 Delivered Stress

The sweep therefore supports robustness to the tested changes in tool mass and centre-of-mass height. It does not establish robustness to unbounded hydraulic variation. A common authority perturbation was used to compare all stacks more directly. This separates a reduction in realised actuator motion from the specific randomisation factors used during deployment.

7.4.2 Actuator Authority

Actuator authority was reduced by scaling the realised actuator displacement:

$$\Delta x_k^{\text{applied}} = \kappa \, \Delta x_k^{\text{actuator}}, \qquad \kappa \in (0, 1]. \tag{7.1}$$

Here, $\kappa = 1$ represents nominal authority and smaller values reduce the motion delivered to the plant. The same perturbation is applied after each stack's actuator representation, which makes the comparison independent of the individual factor parameterisations.

No stack reached the predefined safety or ground-contact limit over the tested authority range. A hard robustness boundary was therefore not observed. The adaptive residual maintained the frozen residual's nominal lateral error down to $\kappa = 0.30$, which indicates an authority margin within AF. IF had the lowest lateral error across the cross-stack comparison. The authority result should therefore be read as an improvement over the frozen adaptive baseline, rather than as a claim of overall superiority.

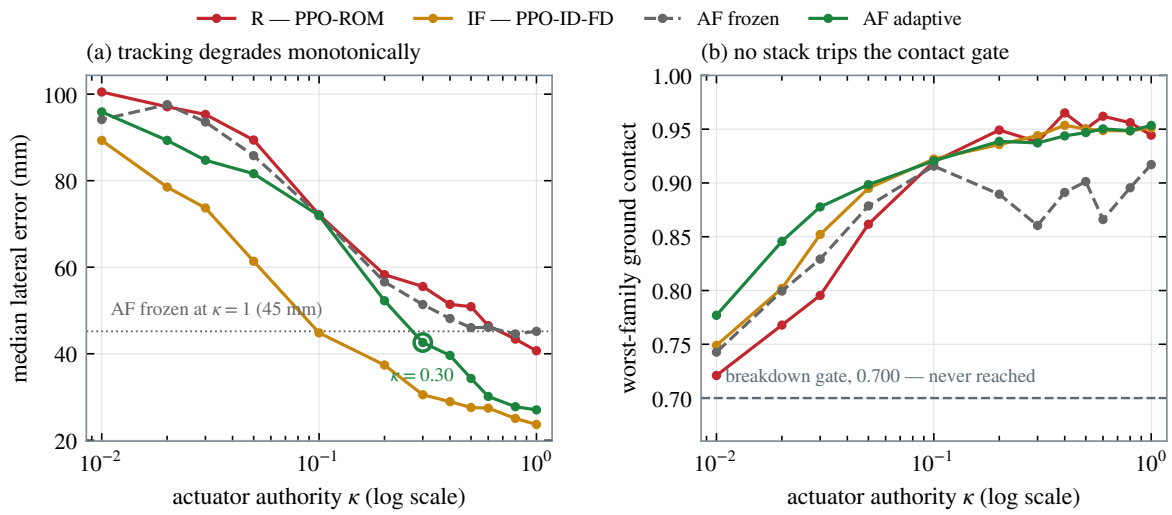

Figure 7.4: Lateral tracking and ground contact as actuator authority is reduced. The frozen and adaptive arms are compared within AF.

Table 7.5: Authority at which median lateral error first reaches 1.5 times its full-authority value.

Stack	nominal error (mm)	threshold (mm)	onset κ
R	40.69	61.04	0.10
IF	23.66	35.49	0.20
AF frozen	45.18	67.78	0.10
AF adaptive	27.03	40.55	0.30

Table 7.6: History-conditioned residual against the frozen base over five training seeds. An asterisk indicates that the paired bootstrap interval excludes zero.

Seed	cut-height gain	lateral gain	interpretation
17	+39.1 %*	+84.9 %*	improvement
23	+35.5 %*	+4.0 %*	improvement
101	+74.7 %*	+5.9 %*	improvement
42	+43.4 %*	−16.1 %*	trade-off
7	+4.5 %	+55.7 %*	lateral improvement

7.5 Online Adaptation

7.5.1 Experiment and Measurement

The online-adaptation experiment changes cutter mass during each episode. The change time and direction vary across environments. The three compared arms share the frozen base policy, residual architecture, training budget, and reward. The base has no residual, the control has a residual with a zeroed history window, and the adaptive arm receives recent state-action history.

Changing mass during an episode creates a condition that cannot be addressed by one fixed response. The varying change time prevents the residual from relying on a fixed event schedule.

Results use 880 paired environments across five held-out meshes. Lateral and cut-height effects are evaluated separately. Cut-height effects smaller than approximately 10 % are treated as unresolved because of the measurement resolution; the corresponding lateral threshold is approximately 1.5 %. The paired design ensures that each arm is assessed under the same terrain, trajectory, and mass-change condition.

7.5.2 Performance Across Seeds

The adaptive residual improves at least one metric in four of five seeds, and no seed degrades both metrics. The size of the improvement varies across seeds, so the result supports a beneficial but variable effect. The trade-off at seed 42 also shows why lateral and cut-height errors are not combined into one score.

The percentages compare paired environments. A positive value means that the adaptive arm has lower error than the frozen base for the median paired environment. The bootstrap marker indicates that the paired difference is resolved by the stated statistical test.

7.5.3 Contribution of History

The capacity-matched control separates the effect of the residual from the effect of the history it receives.

The residual capacity alone accounts for most of the cut-height improvement in four seeds. The effect of live history is mixed across seeds and metrics. Thus, the residual improves on the frozen base, but the full average gain cannot be attributed to history alone. This

Table 7.7: Decomposition of the adaptive gain. The control and adaptive arms have the same capacity; only the adaptive arm receives live history.

	base → control		control → adaptive	
Seed	cut	lateral	cut	lateral
17	+63.9 %*	−1.2 %	−21.1 %*	+80.3 %*
23	+77.2 %*	+42.2 %*	−21.2 %*	−22.8 %*
101	+90.9 %*	+56.2 %*	−5.2 %	−29.9 %*
42	+1.3 %	−62.2 %*	+38.4 %*	+124.9 %*
7	+95.0 %*	+39.7 %*	−46.0 %*	+12.2 %*

Table 7.8: Lateral gain beyond the sham condition. Every interval excludes zero.

Mass-change magnitude	total gain (mm)	gain beyond sham (mm)	interval (mm)
1.0 (sham)	16.557	n/a	n/a
1.5	20.045	3.796	[2.303, 6.105]
2.0	20.487	3.912	[1.693, 6.415]
2.5	20.490	4.073	[2.318, 6.697]
3.0	23.705	7.410	[4.785, 10.503]

comparison is necessary because a larger residual can improve control even when it receives no information about the current plant.

The control arm is not a weaker adaptive arm. It has the same trainable capacity and optimisation conditions. It isolates the effect of live history from the effect of adding residual capacity.

7.5.4 Plant-Dependent Contribution

To isolate plant-dependent adaptation, the mass-change magnitude was varied and compared with a sham condition in which the event occurred without a change in mass. The sham measures the benefit of the residual under the same event schedule but without plant variation. The lateral benefit beyond the sham increased with the mass change.

The additional lateral gain rises from 3.796 mm at a 1.5-fold mass change to 7.410 mm at a 3.0-fold change. This trend is monotonic and resolved at every tested magnitude. It shows that part of the lateral benefit is plant dependent. At the largest change, this component represents about 31 % of the total lateral gain. The corresponding cut-height benefit does not vary consistently with mass-change magnitude and is therefore interpreted as a general residual effect rather than a plant-dependent effect. The experiment therefore separates a general improvement from the smaller component that responds to the change in plant mass.

The result is limited to the perturbation and channel tested here. It does not show that history conditioning improves every plant change or every task metric. It shows that the lateral correction contains a measurable component that scales with the mass change.

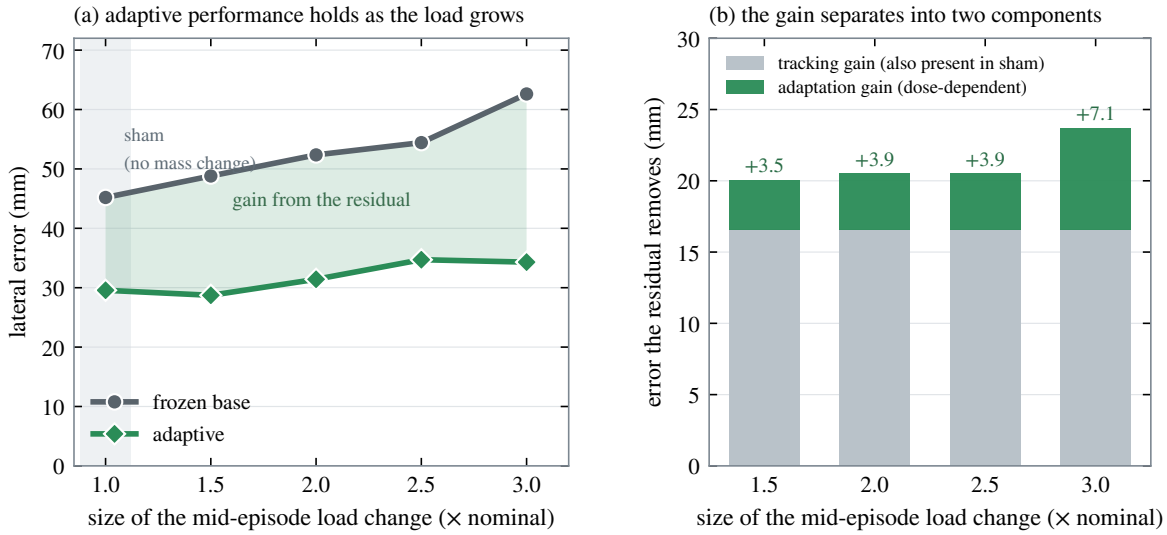

Figure 7.5: Lateral performance across mass-change magnitudes. The sham condition separates general residual improvement from the contribution that increases with the plant change.

7.6 Integrated Discussion

RQ5: Effects of Actuator Representation

The inverse and forward actuator representation improved average tracking in the matched comparison. It also reduced command reversals substantially. Although the hydraulic transport delay was short, both stacks responded to a reference step on a much longer time scale. Command limits and policy dynamics, rather than the hydraulic delay, dominated closed-loop settling.

The tested stacks transferred to held-out terrain, trajectories, and plant configurations. The authority experiment showed graceful degradation but did not identify a hard safety boundary. Within AF, the history-conditioned residual retained the frozen residual's nominal tracking at lower actuator authority. This is an authority-margin result, not evidence that it is the best cross-stack controller.

RQ6: Evidence for Online Adaptation

The history-conditioned residual improved on the frozen base in most training seeds. The capacity-matched comparison shows that residual capacity explains a large part of this result, so the overall gain cannot be assigned to history alone. However, the sham-controlled dose experiment provides direct evidence of a plant-dependent lateral contribution: the benefit increased with the size of the mass change.

Online adaptation is therefore supported for the lateral channel under the tested mass changes. Its contribution depends on plant variation reaching a channel that the correction can influence and on the correction retaining usable command authority. These conditions define the scope of the result.

Relation to the Research Gap

The results show that history-conditioned control should be evaluated separately from added model capacity. A sham-controlled perturbation is useful because it distinguishes a general performance improvement from a response that scales with plant variation. For the hydraulic manipulator studied here, the measurable plant-dependent benefit is limited to lateral tracking under the tested load changes.

8. Conclusion and Future Work

This thesis asked how hydraulic actuator dynamics learned from high-fidelity simulation can be reduced, transferred into a computationally efficient robotic simulation and used for reinforcement-learning control, and under what conditions online adaptation adds anything to the robustness those learned models already provide. The six research questions of Section 1.4 are answered in the order they were asked, which is also the order in which the answers depend on one another.

8.1 Summary of Findings

The findings are stated in the order of the research questions. The first four concern how hydraulic actuation is learned, transferred and controlled, and the last two concern online adaptation.

The inverse and forward representation identified in Simscape transfers to the deployment simulation without retraining (RQ1, RQ2). The inverse model was trained on Simscape data alone and evaluated on all 118 IsaacLab deployment runs, drawn entirely outside that training distribution, where its teacher-forced one-step extension error is millimetre-scale and nearly flat across domains: 2.36 mm on the combined runs against 3.36 mm on nominal. This claim is limited to the one-step regime used by the policy. The inverse model remains frozen downstream. The reduced-order model received a small bounded calibration against the deployment plant.

Predictive accuracy must be accompanied by a control-authority test (RQ1). Input ablation and open-loop excitation showed that a learned actuator must be tested for dependence on the command, not only for prediction error. With the accepted excitation design, removing the valve input increases rollout error by a factor of 1.81. This is a general validation requirement for learned models used in control.

Two actuator representations both support closed-loop control, and they do not produce the same controller (RQ3, RQ5). The command-conditioned reduced-order model reaches $R^2 = 0.998$ on the internal valve state and 0.97 on extension, with a diagonally dominant extension–command Jacobian of condition number 2.14, and it runs inside GPU-parallel training. The inverse–forward chain composes two separately

learned models in series, which could have compounded their errors and did not. The two representations differ in the control strategy they induce from the same policy interface. R spends the least command and reverses about 40 times per second, against 10.6 for IF, which converts greater actuator effort into lower lateral error. The choice of representation therefore affects the resulting controller and is not an implementation detail.

R and IF agree on every recorded training field and differ only in the actuator path, so the difference above *is* attributable to the representation. AF is a separately trained run at a different seed, a lower actor learning rate and no frozen-actor warm-up, reported as a deployed controller rather than a ranked competitor. Only one training run exists per stack, so seed variance is not estimated even for the matched pair, and claims about the adaptation path are carried instead by the paired within-stack comparison of Section 7.5.3.

A learned policy can replace the classical controller at runtime (RQ4). The damped-least-squares controller supplies privileged demonstrations during training and is then removed; the distilled policy emits desired cylinder motion directly into the learned actuator and is the entire deployed controller. Over 33 evaluation cells IF holds a mean lateral error of 70.3 mm, a median of 28.1 mm and 95.29 % ground contact (Table 7.1). The deployed policies remain anchored to the demonstrated behaviour because the reward is an imperfect proxy for the acceptance gates.

The robust stacks generalise without online adaptation, but no robustness boundary was located (RQ2, RQ5). Every closed-loop result is measured on held-out meshes, eleven trajectory families scored individually, and plants drawn from sixteen reserved factor configurations disjoint from training by construction. Under those conditions IF holds the lateral criterion in 32 of 33 cells at a median of 0.0281 m, and the frozen inverse model beneath it, identified on Simscape alone, transfers to all 118 deployment runs with millimetre-scale one-step error, flat across domains. This is the broadest claim the work supports and it requires no plant estimation.

The factor sweep supports the tested mechanical variation, but not unbounded hydraulic variation. On the inverse and forward path, hydraulic factors are not applied; on R they remain inside the identification range (Section 7.4.1). Reducing realised actuator motion to one per cent of nominal did not trigger the safety or ground-contact gates. The result therefore establishes an authority margin, not a hard robustness boundary (Section 7.4.2).

The remaining findings concern adaptation itself (RQ6).

Plant identity requires suitable excitation (RQ6). In the task-control data, command and state are nearly redundant. Open-loop excitation makes plant identity observable and is therefore used for the model-level results. This finding concerns the present data collection and does not exclude dedicated closed-loop identification methods.

The estimator is not the limiting component (H1, H2, RQ6). Under the chain gate of Section 5.6.3 the causal estimate is not worse than the privileged latent read from the true factors, so whatever the chain loses, it does not lose it at estimation. What the estimate carries is mainly seal condition (Section 5.5.1), which acts through pressure, and pressure is the one channel where the latent's advantage is measurable. That is also the channel the controller does not integrate, which is the first sign of the delivery problem below.

Online adaptation requires usable correction authority. The deployed residual acts directly on the valve command. It improves at least one metric in four of five seeds, and no seed degrades both metrics. The strongest cut-height gain is 74.7 % (Section 7.5).

A capacity-matched control shows that residual capacity explains much of the cut-height gain. The contribution of live history is mixed across seeds and metrics, so the overall gain cannot be attributed to sensing alone (H4).

The authority axis gives the same pairing a second and cleaner reading. Against its own frozen arm the adaptive arm holds the frozen arm's full-authority tracking down to 30 % of nominal actuator authority, a 3.3-fold margin, and leads it across the useful range of κ (Section 7.4.2). This is a within-stack statement: IF remains the stronger controller on that axis at every rung, and the result is a margin rather than an extended envelope, since no breakdown boundary was reached.

The sham-controlled dose experiment isolates a lateral component that increases monotonically with the size of the mass change. This provides direct evidence of a plant-dependent contribution in addition to the general residual effect.

Plant-dependent adaptation contributes when the variation reaches a channel the correction can influence, the correction has authority, and the data contains information about the plant. In this work, the resolved effect is limited to lateral tracking under the tested mass changes.

8.2 Design Implications

The results carry design consequences for each of the three parts of this work: how a hydraulic actuator is identified and transferred, how a policy is trained to act through it, and how plant adaptation is fitted around both. They are stated below in that order, because each stage constrains the next.

Identification and Transfer

Design the excitation for identifiability, and test the model by input ablation. Task-control data can make state and command nearly redundant. Dedicated excitation makes command effects identifiable. A learned actuator should therefore be accepted on both prediction accuracy and loss of performance when the command input is removed.

Embed known physics and bound learned corrections. The reduced-order model propagates pressure, force, velocity, and extension through explicit updates. Its learned terms are bounded, preventing unphysical state changes inside the training loop. This structure also represents the pressure dependence of valve authority directly.

Prefer a frozen model with a bounded, gated correction to retraining. The inverse model transferred across the deployment plant family at millimetre scale in every domain, from 2.36 mm to 3.36 mm one-step, and with an ordering that does not follow plant severity: the nominal plant scores worst and the combined draws best (Table 5.2). It was therefore deployed frozen everywhere downstream. Where a deployment offset did have to be absorbed, it was absorbed by a correction capped at 3 mm and gated to vanish at zero command, with every hydraulic parameter frozen and the original data retained as replay. That construction preserves the identified physics, keeps a commanded hold a hold,

and confines the deployment-specific part of the model to a component whose authority is bounded by construction.

Control Methodology

Constrain policy updates when the reward is not the acceptance criterion. The physical acceptance gates and the reward are not identical. A forward-KL anchor to a verified reference policy bounds policy drift, and the anchor advances only through checkpoints that satisfy the gates. Reward terms should represent quantities that are also treated as acceptance constraints.

Where no baseline is inherited, the action contract is the safety argument. The deployed policy is not a residual: a zero action leaves the arm stationary, and nothing is summed with the policy output at runtime. Authority is therefore bounded by the chain the action passes through rather than by proximity to a nominal controller, with a rate limit, a filter, a shield that fades authority as travel headroom is consumed, a command scale, and an anti-windup cap on how far the integrated target may lead the measured extension. That last cap is load-bearing and axis-dependent: relaxing it from 18 mm to 50 mm halved lateral error on the families needing a sustained standing offset, while the same relaxation on the head-pivot cylinder drove alignment from 0.9507 to 0.4636.

Plant Adaptation

Measure identifiability and correction authority before training an estimator. The identifiability envelope determines whether plant identity is present in the data. The delivered-stress check determines whether a perturbation reaches the plant and whether a correction can influence the command. Both tests should precede estimator training.

Where the estimate is used also matters. A dynamics model can integrate a changing estimate over time, whereas a memoryless policy receives that change as input variation. Filtering on the policy path may therefore be useful, but its appropriate time constant remains to be determined on the deployed stack.

8.3 Contributions

Section 1.4 stated six contributions. Their standing at the end of the work is as follows, with the supporting evidence as summarised above rather than repeated.

Transferable hydraulic actuator learning is delivered: both representations were identified from Simscape and realised in IsaacLab, and the inverse model transfers without deployment-specific retraining. The **reduced-order actuator model** is delivered and carries everything downstream, being accurate enough to serve as the deployment plant and cheap enough to evaluate inside GPU-parallel training, which the stiff implicit solver of the model it replaces cannot support; it was calibrated against IsaacLab and is therefore the reference condition rather than an independent competitor.

The **control-relevant identification and validation** procedure is delivered and is the contribution most likely to transfer independently of this machine, covering both the input-ablation test that exposed a forward model with no command authority and the identifiability audit that established the excitation envelope.

Policy-only control through learned hydraulics is demonstrated rather than certified. On worst-mesh, worst-family statistics both stacks fail the lateral gate and IF fails head alignment badly (Section 7.2), so the evidence supports that a policy can drive the task through learned hydraulics with useful tracking, not that either stack satisfies the whole eight-gate acceptance contract.

The **decision procedure for online adaptation** is delivered, and its main product is the discipline that produced it rather than any single number. Two further findings exist only because of that discipline: that a selection window covering 8.9 of 26 m certifies a third of an episode, and that a quantity which is gated but not priced is traded away by the optimiser.

Online adaptation on a saturating actuator stands as a demonstrated mechanism. The median of five seeds does not meet the pre-registered bar, so it is not yet a general method.

Two enablers are recorded rather than claimed. The causal plant estimator and the bounded adaptation path behave as specified, with no privileged access, no future information, bounded output and reversion to the robust actuator path on invalid input. Of the eight-gate promotion criterion, one property is worth reuse: five gates separate the matched pair and carry the comparative claims, while lateral, hard impacts and tool acceleration do not (Section 3.9.1).

8.4 Limitations

The evidence is simulation-to-simulation, and the deployment plant is itself a learned model. No claim here is a claim about hardware.

The differential pressure the estimator consumes is produced by the actuator model rather than measured, so transferring the model-level result to hardware requires instrumentation the FME 500 does not carry. An observation contract remains unresolved alongside this: every result uses the pressure-inclusive row, while the stricter deployable alternative excludes measured differential pressure. Since the estimator's advantage is concentrated in that channel, the deployable variant should be expected to perform worse, and by how much is not established.

The closed-loop comparisons resolve differences to a few per cent over 33 paired conditions and do not establish equivalence. The available plant variation bounds the disturbance rather than the architecture, and more tightly than the sweep's labels suggest: a stack shown to absorb a 40% reduction in valve effectiveness and a nominal doubling of every deployment factor has been shown to tolerate what actually reached it, which on the inverse/forward plant is a mechanical perturbation only. The three stacks are matched on evaluation but not on training budget, so the comparison describes three deployed controllers rather than three equally tuned training procedures.

Policy-training variation is not quantified anywhere in this work, and it bounds two claims differently. R and IF are each represented by one training run, so the 33 paired cells measure terrain and trajectory variation and say nothing about how far a second run of either would move; the actuator-representation result rests on that assumption. The adaptation result is quantified across seeds and bounded by them: two of five improve both errors, one trades, two are worse, so the pre-registered median is not met. The

defensible claim is a plant-dependent lateral contribution under the tested conditions, not an improvement on both errors for every training seed.

Finally, the bound is a bound on this loop. A commanded step cannot complete faster than 417 ms at the command rate limit while the plant answers in 117 ms, so the task response is set downstream of anything a latent could change within a transient. That is close to the least favourable case for adaptation understood as *response to an event*, and it is consistent with the measurement: the adaptive arm's advantage does not widen when the load moves. It says nothing against adaptation understood as continuous conditioning on the operating point, which is the form that pays here. The finding generalises through the mechanism rather than through the number.

8.5 Future Work

The evidence in this thesis is simulation-to-simulation. The next step is to evaluate the learned actuator models and distilled policy on the target machine.

That step begins with instrumentation, because the FME 500 as delivered carries no cylinder-chamber pressure transducers, and the differential pressure the models propagate exists here only as an actuator-model state. Fitting pressure transducers, extension sensing and a rod load pin converts three modelled quantities into measured ones. The collection instrument is already designed and validated: the open-loop excitation schedule of Chapter 4 was built to make the valve causally identifiable, and the identifiability audit of Section 4.7.5 applies to machine data unchanged.

Hardware data would support a bounded correction to the frozen inverse model, with hydraulic parameters retained and the original dataset used for replay. Input ablation should verify that the correction remains dependent on the valve command. The reference-anchored policy provides a suitable starting point because it limits policy drift from a verified checkpoint.

A. System Parameters and Coordinate Mappings

The kinematic structure, the cylinder-to-joint mapping and the calibrated cutter Jacobian are derived in Section 3.2 and are not repeated here. This appendix records the quantities a reader needs in order to reproduce a run and cannot read off those derivations: the rate of every layer, the limits the command path enforces, and the geometry of the scored course. Every value below is read from the shipped configuration or from a stored evaluation record, and the source is named in each case.

A.1 Rates

Five different rates operate in this system and confusing them is the most common way to misread a result. The control step is 1/60 s, the physics substep is half that because the environment runs with a decimation of two, and every learned component operates on the 20 Hz grid at which the reduced-order model and the identification data were sampled.

Table A.1: The rate of every layer. Control and physics rates are read from the stored evaluation records (`control_dt_s`); the learned-model sample time is stored inside the frozen forward-model checkpoint (`sample_time_s`); the AAM update period is the runtime configuration `update_period_ticks`.

Layer	Rate	Period
PhysX physics substep	120 Hz	8.33 ms
Policy and control step	60 Hz	16.67 ms
Inverse / forward model	20 Hz	50 ms
AAM inference update	20 Hz	50 ms (every 3rd control step)
Simscape logging grid	20 Hz	50 ms

Two derived latencies follow from Table A.1 and are used in Section 7.3. Both belong to the offline AAM and the ID adapter it fed, which is the earlier path of Section 5.6.3. The AAM consumes a history of 20 rows on the 50 ms grid, so its input window is exactly 1.0 s

and no estimate at all is available for the first 1.0 s after a reset; until then that adapter is bypassed and the frozen inverse action is returned unchanged. Its output smoother averages five consecutive predictions, which adds a further 250 ms of group delay.

The deployed AF residual of Section 6.7.3 has neither latency and no inverse model to fall back to. Its encoder reads 20 rows on the control grid, so its window is 333 ms, the window is carried in the observation and zero-filled at start rather than gated, and no smoother is applied. Nothing is bypassed: the residual is exactly zero at construction because its head is zero-initialised, so a zero-filled window reproduces the frozen base rather than switching to another path.

A.2 Command-Path Limits

The command path enforces four limits, all of which bound the policy rather than advise it. They are listed in Table A.2.

Table A.2: Limits enforced on the command path, per cylinder where applicable. Valve command limits and the adapter residual bound are stored inside the frozen ID-adapter checkpoint (`command_limits`, `residual_bound`); lead caps and the c3 reserve are read from the stored evaluation records. The residual bound row belongs to the ID adapter of Section 5.6.3. The deployed AF residual is bounded differently, by $\alpha \tanh(\cdot)$ at $\alpha = 0.5$ on the valve command (Section 6.7.3).

Limit	c1	c2	c3	c4	Meaning
Valve command limit	0.95	0.95	0.80	0.80	accumulated command is clipped here
Adapter residual bound	0.15	0.15	0.15	0.15	maximum adapter correction
Target lead cap (m)	0.03	0.03	0.03	0.02	how far the target may lead the state
c3 extension reserve	n/a	0.008 m on c3			solver-enforced stroke reserve

The lead cap is the mechanism by which the reference governor prevents the commanded target from running away from the measured state, and the c3 reserve is enforced by the solver while headroom is still scored against the original stroke, so the gate is not satisfied by moving the goalposts. The ID adapter's residual bound is what made that path fail safe: an adapter output that is invalid, uninitialised or late is replaced by zero, which returns the stack to the frozen-inverse configuration rather than to an untested one. The deployed residual fails safe by the same principle at a different site. Its correction is bounded by $\alpha \tanh(\cdot)$ whatever the head produces, and setting α to zero at evaluation time recovers the frozen base exactly, which is the property the control arm of Section 7.5.3 relies on.

A.3 Course and Episode Geometry

Every closed-loop result in Chapter 7 uses the geometry in Table A.3. The scored window is the full course: the count of control steps differs between evaluations only because the carrier speed does, and it is the travelled distance rather than the step count that defines the window.

The right-hand column is included because stored records under it exist in the project and are not comparable to the results of Chapter 7; Section 7.1.3 explains why one such record, which passes all eight gates, is not used.

Table A.3: Course geometry and the two evaluation speeds that appear in this thesis. Values are read from the stored evaluation records (`tracking_active_length_m`, `max_truck_travel_m`, `steps`).

Quantity	Chapter 7 protocol	Earlier 2 km/h records
Tracking line length	26.0 m	26.0 m
Carrier speed	4.0 km/h	2.0 km/h
Control steps scored	1403	2808
Truck travel achieved	25.98 m	26.00 m
Episode length	47.3 s	46.8 s
Trajectory families	11	8

B. Dataset Manifests and Factor Ranges

B.1 Trajectories

Each run receives its own reference trajectory, sampled rather than scripted. Between twelve and twenty waypoints are drawn by Latin hypercube sampling within the joint limits, and the arm moves between them without dwelling, so that the data contains transitions rather than long static holds. Segment durations are drawn from four speed classes, slow (1.5–2.5 s), medium (0.8–1.5 s), fast (0.3–0.7 s), and variable (0.5–2.5 s), and about a third of the runs receive additional exploratory perturbations on top of the waypoint path.

The trajectory design limits long static intervals so that the collection contains transitions as well as steady motion. This improves its usefulness for learning actuator dynamics.

B.2 Execution and Logging

Each run simulates 22 s of operation. The solver is the stiff implicit `ode23t` integrator in accelerator mode with relative and absolute tolerances of 10^{-2} and 10^{-3}. The loose tolerances are a deliberate throughput compromise, and their effect is bounded by the consistency checks of Section 4.10. Output is written on a fixed 0.05 s grid, giving 441 logged samples per run. The batch is checkpointed to disk every hundred runs.

For every run the archive stores the time vector, the five valve commands, the four cylinder extensions, the four chamber pressure differences, the four rod forces, the five joint positions and velocities, the reference position and velocity signals, and the complete parameter structure that produced the run. That last item is what makes the batch usable for the latent-adaptation work: the privileged plant descriptor $\boldsymbol{\xi}$ of Section 2.6 is not something that has to be reconstructed later, it is recorded alongside the trajectory.

The cost quantifies the argument made in Section 2.1.3. A single 64 s open-loop run takes roughly 2000 s of wall time on the collection machine, a real-time factor of about 0.03. The Simscape plant runs some thirty times slower than the machine it models. No policy-optimisation loop that needs millions of environment steps can be built on that, which is why the learned actuator surrogates exist.

B.3 Parameter Ranges and Dataset Audit

Table B.1: Randomised hydraulic parameters. The seal-wear coordinate is drawn independently for each of the four cylinders and drives three physically linked quantities at once.

Parameter	Nominal	Randomised range	Acts on
Entrained air fraction ε_{air}	0.005	0.001 – 0.080	fluid compliance
Dead-volume fraction ε_{vol}	0.020	0.005 – 0.080	fluid compliance
Discharge coefficient c_d	0.70	0.56 – 0.84	valve flow gain
Seal wear $w_{\text{seal},i}$	0	0 – 0.25 (per cylinder)	clearance, leakage, spool lag
Relief setting p_{set}	27.6 MPa	24.8 MPa – 30.4 MPa	pressure ceiling

Table B.2: Randomised mechanical parameters. The rigid-body perturbations are applied independently to each of the fourteen bodies.

Parameter	Nominal	Randomised range
Body mass	CAD value	±20 %
Centre-of-mass offset	CAD value	±5 %
Inertia tensor and products	CAD value	consistent with mass draw
Joint stiffness, J1 and J2	60 kN/m	40 – 120 kN/m
Joint stiffness, J3 and J4	60 kN/m	40 – 120 kN/m
Joint stiffness, J5	25 kN/m	15 – 60 kN/m
Joint damping	250	100 – 600

Table B.3: The six randomised physical factors of the IsaacLab deployment collections and their final ranges. One factor vector is drawn per run and held fixed for the whole run.

Coordinate	Range	Domain	Acts on
e_{flow}	$[-0.25, 0.25]$	hydraulic	valve flow gain
$e_{\text{compliance}}$	$[-0.25, 0.10]$	hydraulic	effective chamber compliance
e_{seal}	$[0, 1]$	hydraulic	seal wear, leakage, spool lag
e_{pressure}	$[-0.15, 0.15]$	hydraulic	supply/relief pressure level
e_{mass}	$[0.8, 1.5]$	mechanical	cutter mass scale
e_{COM_z}	$[-0.10, 0.05]$	mechanical	cutter centre-of-mass height

B.4 The Frozen Deployment Collections

Two collections are recorded in the deployment plant and both are frozen: the task-controller collection and the open-loop excitation collection of Section 4.7.5. They share their factor design, domain composition and run-level split, and differ only in how the arm is driven. Each manifest records 118 runs, one factor vector per run, held fixed for the duration of that run. The model-level results in this thesis are fitted and scored on the open-loop collection, for the identifiability reason given in Section 4.7.5. The composition common to both is given in Table B.5.

Table B.4: Audit of the strict open-loop excitation export. All checks are computed by the extraction stage and stored with the data; the export is accepted only if every one of them passes.

Check	Result	Interpretation
Rows, runs, columns	25 600 / 4 / 49	four complete runs
Median sample time	0.01 s	as commanded
Maximum sample-time error	9.1×10^{-15} s	numerically exact grid
Maximum Δextension error	0	targets match state differences
Maximum Δvelocity error	0	targets match state differences
Maximum $\Delta(\Delta p)$ error	0 Pa	targets match state differences
Non-finite values	0	no solver escapes
Valid fraction	0.937	startup, gaps, shutdown excluded
Command rank	4	all four channels independent
Lagged excitation rank	80	history matrix full rank
Valve saturation	8.5 % per channel	limits reached but not dominant
Stroke coverage	12 % – 33 %	above floor, below target

Table B.5: Composition of the frozen deployment collections, counted directly from the manifest. The run-level split is 102 training runs against 16 held-out test runs; the training portion is further divided 82/20 into fitting and validation runs, so that no window from a validation or test run is ever seen during fitting.

Domain	Runs	Factors varied
`combined`	64	all six
`hydraulic`	40	flow, compliance, seal, pressure
`mechanical`	12	mass scale, centre-of-mass height
`nominal`	2	none (factors at nominal)
Total	118	
`train` split	102	further split 82 fit / 20 validation
`test` split	16	held out

The split is at *run* level rather than window level, and this is the single most important property of the collection. Within one run the factor vector is constant, so windows drawn from the same run are highly redundant; a window-level split would place near-duplicates of the same plant on both sides of the boundary and every held-out score in this thesis would be optimistic. The manifest stores the factor key of every run alongside its split label, so this property is checkable rather than asserted.

C. Model Architectures, Normalisers and Checkpoint Hashes

This appendix pins the frozen artifacts. Every model-level and closed-loop result in this thesis was produced with the exact files listed in Table C.1, and the SHA-256 digests are given so that a later reader can establish whether a file on disk is the one that produced a number here. The adaptation results in Section 7.5 use the frozen base `phase1_init_70.pt` and the residual checkpoints listed below.

Table C.1: Frozen artifacts and their SHA-256 digests (first 32 hex digits shown). "Frozen" means the file's weights were not updated by any subsequent stage. The residual checkpoints are the trained arms of Tables 7.7 and 7.6; each is evaluated on top of the frozen base above it.

Artifact	Role	SHA-256 (truncated)
`aam.pt`	AAM, frozen	`cd6a7f0b118b96c7ce66b889b9b3e422`
`zhat_adapter.pt`	ID adapter, frozen	`30629a4f83db5c2bed3c777bcf82dc57`
`best_id_cylinder_control.pt`	inverse model, frozen	`8e26c52e82d89efe193e001d0a8b07a6`
`best_fd_cylinder_strict.pt`	forward model, frozen	`1f41e44309a8f320430a3382ca88ce5f`
`rom_s6_aligned/model_598.pt`	R policy	`ce7e3ab9649b1193f123f9f91107bbf8`
`idfd_s6_aligned/model_598.pt`	IF policy	`06533b8e8614b0ad982dd8c238a17772`
`distilled_v1/phase1_init_70.pt`	AF frozen base	`641e476cfcd765a8d68a17980d6cbbf2`
`rebal17_history/model_399.pt`	residual, seed 17	`ea39150fd2919ce96aa1acdbbdc9d65d`
`rebal23_history/model_399.pt`	residual, seed 23	`63da26e9b9710b2f54715d7d70e82d59`
`rebal101_history/model_399.pt`	residual, seed 101	`f5a400f3b7e66b46a0cca11142a67d70`
`rebal42_history/model_399.pt`	residual, seed 42	`5cafe571d429ea8181989ad87c2ad558`
`rebal7_history/model_399.pt`	residual, seed 7	`4e5c3041987c20685e928ac0d5fac192`
`rebal17_nohistory/model_399.pt`	zero-history control, seed 17	`c8b001b27ded55f2f32d5fe47d022a0a`

C.1 The Inverse and Forward Models

Both learned actuator models share one architecture: a linear input projection to 512 units, four pre-norm residual blocks of two 512-wide linear layers each, a final normalisation and a linear output projection. They differ only in their input and output layouts, given in Table C.2. Each carries just over 2.12 million parameters. Both are trained on the

Table C.2: Input and output layouts of the two frozen actuator models, read from the `input_layout` and `output_layout` fields stored in the checkpoints themselves.

	Inverse model	Forward model (strict)
Input width	28	20
Input layout	extension(4), velocity(4), $\Delta p(4)$, force(4), target Δextension(4), target Δvelocity(4), $\Delta t(4)$	valve(4), extension(4), velocity(4), $\Delta p(4)$, force(4)
Output width	4	12
Output layout	valve command(4)	Δextension(4), Δvelocity(4), $\Delta(\Delta p)(4)$
Target mode	normalised Δ valve command	normalised Δ cylinder state
Parameters	2 124 292	2 124 300

Simscape campaign, both operate on the 50 ms grid, and both are frozen for every result in Chapters 6 and 7.

Only the predicted extension increment of the forward model is integrated in the control loop. Its velocity output is retained as a diagnostic.

C.2 The Adaptation Module

This section describes the offline AAM, the encoder trained against a latent target in Chapter 5. The residual encoder deployed in AF is a different component, trained jointly with its head under the reinforcement objective and against no target (Section 6.7.5).

The AAM is a causal one-dimensional convolutional network over the last 20 rows of the 50 ms deployment observable. The input row is 18 columns wide,

$$h_t = \Big[\, q(5),\ \dot{q}(5),\ \Delta p(4),\ u_{t-1}(4) \,\Big],$$

and the trunk is four convolutional stages of width 64 and kernel three, the first carrying a 1×1 skip projection from 18 to 64 channels, each stage normalised. A two-layer head maps the final 64-channel activation to the 2-dimensional latent $\hat{z}$. The whole module is 42 562 parameters, which is what makes 20 Hz inference affordable alongside 256 parallel environments.

The latent dimension is 2 because the module regresses onto the frozen state predictor's latent, which is 2-dimensional by construction. This is the structural reason the deployed latent decodes only one of the six physical factors, and Section 5.5.1 treats it as a finding rather than a configuration choice.

The ID adapter is deliberately small: a two-layer network of width 64, mapping a 30-dimensional feature row to a 4-dimensional residual, 76 stored tensors' worth of normalisation statistics included. Its feature row is the inverse model's 28 inputs concatenated with the 2 latent components, which is the pressure-inclusive contract discussed in Section 7.1.2. Its output is bounded at 0.15 per cylinder (Table A.2) and scaled by a declared residual factor of 0.25. That factor applies to the ID-adapter results, which are the null result of Section 5.6.3 and the latent arms of Appendix E. It is not the deployed residual's authority.

The adapter checkpoint additionally stores the digests of the inverse and forward models it was fitted against, so a latent adapter cannot silently be paired with a different frozen actuator than the one it was trained for.

C.3 The Deployed Residual

The residual of Section 6.7.3 is the component every adaptation result in Section 7.5 is produced by, and it shares nothing with the ID adapter above except the idea of a bounded correction. It acts on the valve command, not on the inverse model's output, and there is no inverse model in the stack for it to correct.

Its encoder ϕ reads 20 rows of (base observation, previous action) on the 60 Hz control grid, a 333 ms window carried in the observation vector rather than in a buffer held inside the module. The head g_ψ is a small multilayer perceptron with a zero-initialised final layer, producing a per-cylinder correction bounded by $\alpha \tanh(\cdot)$ at $\alpha = 0.5$. Together they hold 53 836 trainable parameters against the 59 588-parameter frozen base, and the latent width is 8.

Every arm reported in Section 7.5 was trained with the same command, differing only in the training seed and, for the control arm, in the flag that zeroes the history window: 256 environments, 400 iterations, history 20, latent dimension 8, $\alpha = 0.5$, learning rate 3×10^{-4}, and the rebalanced lateral kernel at standard deviation 0.09 m and weight 4.0. The evaluation plant is the mid-episode load change: bimodal cutter mass at 1.0 and 3.0, step time 20 s, jitter 8 s. Checkpoints are listed in Table C.1; every arm is scored at iteration 399, its last, under the rule declared before the campaign.

D. PPO and IsaacLab Configuration

D.1 The Observation Vector

The policy observation is 66 columns for every stack. AF deploys this same actor frozen and unmodified; its adaptation residual reads a history window carried alongside the observation rather than appended to the actor's input (Section 6.7.3). The 66–68 row below records the two latent columns of the earlier latent-conditioned design, retained for reference. Table D.1 gives the layout by index. The final twelve columns are the model-error block, and they exist because the policy is a memoryless multilayer perceptron: it already observes the actuator's prediction and the current measured extension, but forming the error between them requires differencing across two steps, which a feedforward network cannot do. Without these columns the standing model error of a learned actuator is invisible to the policy no matter how long it trains.

The truck-speed channel is normalised using statistics collected over the full deployment speed range. This prevents a fixed-speed normaliser from distorting the observation at other speeds.

D.2 The Action

The action is four-dimensional in every stack, and it does not mean the same quantity in all three. Under R and IF it is a desired cylinder motion, and what converts it to a valve command is the actuator under test, which is the only free variable in that pair. Under AF the same four-dimensional bounded action is read as the valve command itself, scaled to span the admissible valve range, because the inverse model was removed from the command path (Section 6.7). That difference is why no table ranks AF against the other two on task error. The limits enforced between the action and the plant are given in Table A.2.

D.3 The Reward

The ten reward terms and their shaping functions are given in Table 6.2 and are not repeated here. Two weights behind that table are the ones a reproduction is most likely to

Table D.1: Policy observation layout. Indices are zero-based and half-open. Columns 0–53 are the task block, 54–65 the model-error block, and 66–67 belong only to the earlier latent-conditioned design.

Index	Width	Contents
0–9	9	verge height profile ahead of the cutter
9	1	cut-height error
10	1	lateral error
11	1	cutter vertical velocity
12–16	4	measured cylinder extensions
16–20	4	measured cylinder velocities
20–24	4	cylinder headroom, low side (normalised)
24–28	4	cylinder headroom, high side (normalised)
28–32	4	previous applied action
32–36	4	current valve commands
36–40	4	actuator-predicted Δextension
40–44	4	commanded target lead
44	1	head-down alignment
45	1	head-down alignment rate
46	1	truck speed
47	1	scan valid fraction
48–50	2	world-XY components of the cutter's local $+Z$ axis
50	1	normalised cutter–terrain interaction force
51–54	3	path preview at 0.5, 1.0, 2.0 m ahead
54–58	4	cylinder tracking residual (measured − commanded)
58–62	4	actuator prediction residual
62–66	4	normalised cylinder effort
66–68	2	$\mathbf{z}_{\mathrm{AAM}}$, earlier latent-conditioned design only

get wrong: contact loss is weighted at 20.0 and hard impacts at 20.0, against a support reward of 4.0. The hard-impact weight was introduced after the measurement described in Section 3.9.1: the quantity was gated but not priced, and the optimiser traded it away while its return rose.

Reward Design Rules

The lateral-error term uses bounded shaping so that its gradient remains useful far from the reference line. Worst-case weighting is applied only to non-positive reward terms. The reward logs every term separately, allowing its contribution to be checked during training.

D.4 PPO Hyperparameters

Table D.2 lists the optimisation settings of the three evaluated runs, read from their stored run manifests rather than from the scripts that launched them. The manifests share a schema, which is what makes the comparison in Section 7.1 checkable; the fields on which the runs differ are exactly those discussed there.

Table D.2: PPO configuration of the three policy runs, from `run_manifest.json`. "n/a" marks a field the manifest does not record. Differences between the adaptive run and the matched pair are discussed in Section 7.1. The AF column is the earlier latent-conditioned run `headroom_priced_v1`, which carries the checkpoint selection of Section 7.1.3 and the latent arms of Appendix E; its residual scale of 0.25 is that run's value. The deployed residual is trained separately on a frozen base and its configuration is in Appendix C.3.

	R	IF	AF
manifest schema		`terrain_shared_v2`	
seed	17	17	41
terrain base seed	17	17	41
environments	256	256	256
curriculum stage	6	6	6
actor learning rate	3×10^{-4}	3×10^{-4}	3×10^{-6}
entropy coefficient	0.001	0.001	0.001
frozen-actor warm-up	150	150	0
observation normaliser	freeze	freeze	freeze
observation width	66	66	66
discount γ	n/a	n/a	0.97
GAE λ	n/a	n/a	0.95
steps per environment	n/a	n/a	32
episode length	n/a	n/a	47.3 s
truck speed	n/a	n/a	2–7 km/h, per episode
residual scale	n/a	n/a	0.25
trajectory families		all 11	
actor parameters	n/a	n/a	59 080
critic parameters	n/a	n/a	58 881

The dashes are informative and are left visible rather than filled in from the launch scripts. The two reference runs' manifests do not record their truck speed schedule, discount or

GAE settings, so those cannot be verified from the frozen record; the adaptive run's manifest does record them. Section 7.1 states the consequence.

D.5 The Three R Acceptance Records

Historical acceptance records use different checkpoints, speeds, or scoring horizons. They are not used for the comparison in Chapter 7. That comparison evaluates one checkpoint per stack with a common full-path protocol, held-out meshes, and the same eleven trajectory families.

E. Supplementary Results

E.1 Reproduction

Every closed-loop number in Chapter 7 is produced by the experiment scripts stored with the project under `experiments/E1` through `experiments/E6`, each of which records the checkpoint, the mesh set, the family set, the speed and the scored horizon in its own manifest. The adaptation arms are trained and scored by `PPO_AAM_FD/experiment_scripts`, and the paired statistics of Section 7.5 are recomputed from the stored per-environment reports under `outputs/rif_delay` by `experiments/E3/code/seed_table.py`. The model-level results come from the frozen collection of Appendix B.4.

Every table in this thesis is reproducible from the frozen artifacts, Simscape collections, and IsaacLab collections listed above. Here, reproduction means recomputing reported results from the stored data and checkpoints.

Bibliography

[Åström 95] K. J. Åström, B. Wittenmark, *Adaptive Control*, 2nd ed., Reading, MA: Addison-Wesley, 1995.

[Ba 16] J. L. Ba, J. R. Kiros, G. E. Hinton, "Layer Normalization", arXiv preprint arXiv:1607.06450, 2016.

[Bengio 15] S. Bengio, O. Vinyals, N. Jaitly, N. Shazeer, "Scheduled Sampling for Sequence Prediction with Recurrent Neural Networks", in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 28. 2015, pp. 1171–1179.

[Benner 15] P. Benner, S. Gugercin, K. Willcox, "A Survey of Projection-Based Model Reduction Methods for Parametric Dynamical Systems", *SIAM Review*, vol. 57, no. 4, pp. 483–531, 2015.

[Cho 14] K. Cho, B. van Merriënboer, C. Gulcehre, D. Bahdanau, F. Bougares, H. Schwenk, Y. Bengio, "Learning Phrase Representations using RNN Encoder–Decoder for Statistical Machine Translation", in *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*. 2014, pp. 1724–1734.

[Chua 18] K. Chua, R. Calandra, R. McAllister, S. Levine, "Deep Reinforcement Learning in a Handful of Trials using Probabilistic Dynamics Models", in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 31. 2018.

[Egli 20] P. Egli, M. Hutter, "Towards RL-Based Hydraulic Excavator Automation", in *Proceedings of the IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. 2020, pp. 2692–2697.

[Egli 22] P. Egli, M. Hutter, "A General Approach for the Automation of Hydraulic Excavator Arms Using Reinforcement Learning", *IEEE Robotics and Automation Letters*, vol. 7, no. 2, pp. 5679–5686, 2022.

[Elfwing 18] S. Elfwing, E. Uchibe, K. Doya, "Sigmoid-Weighted Linear Units for Neural Network Function Approximation in Reinforcement Learning", *Neural Networks*, vol. 107, pp. 3–11, 2018.

[He 16] K. He, X. Zhang, S. Ren, J. Sun, "Deep Residual Learning for Image Recognition", in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. 2016, pp. 770–778.

[Hochreiter 97] S. Hochreiter, J. Schmidhuber, "Long Short-Term Memory", *Neural Computation*, vol. 9, no. 8, pp. 1735–1780, 1997.

[Hu 26] P. Hu, T. Wen, D. Zhang, H. Chen, J. Gong, "Bayesian reinforcement learning for adaptive control of energy recuperation in hydraulic excavator arms", *Scientific Reports*, vol. 16, p. 6195, 2026.

[Huber 64] P. J. Huber, "Robust Estimation of a Location Parameter", *The Annals of Mathematical Statistics*, vol. 35, no. 1, pp. 73–101, 1964.

[Hwangbo 19] J. Hwangbo, J. Lee, A. Dosovitskiy, D. Bellicoso, V. Tsounis, V. Koltun, M. Hutter, "Learning Agile and Dynamic Motor Skills for Legged Robots", *Science Robotics*, vol. 4, no. 26, p. eaau5872, 2019.

[Isermann 11] R. Isermann, M. Münchhof, *Identification of Dynamic Systems: An Introduction with Applications*, Berlin, Heidelberg: Springer, 2011.

[Jelali 03] M. Jelali, A. Kroll, *Hydraulic Servo-systems: Modelling, Identification and Control*, London: Springer, 2003.

[Ji 22] G. Ji, J. Mun, H. Kim, J. Hwangbo, "Concurrent Training of a Control Policy and a State Estimator for Dynamic and Robust Legged Locomotion", *IEEE Robotics and Automation Letters*, vol. 7, no. 2, pp. 4630–4637, 2022.

[Johannink 19] T. Johannink, S. Bahl, A. Nair, J. Luo, A. Kumar, M. Loskyll, J. A. Ojea, E. Solowjow, S. Levine, "Residual Reinforcement Learning for Robot Control", in *Proceedings of the IEEE International Conference on Robotics and Automation (ICRA)*. 2019, pp. 6023–6029.

[Kakade 02] S. Kakade, J. Langford, "Approximately Optimal Approximate Reinforcement Learning", in *Proceedings of the 19th International Conference on Machine Learning (ICML)*. 2002, pp. 267–274.

[Karniadakis 21] G. E. Karniadakis, I. G. Kevrekidis, L. Lu, P. Perdikaris, S. Wang, L. Yang, "Physics-Informed Machine Learning", *Nature Reviews Physics*, vol. 3, no. 6, pp. 422–440, 2021.

[Kingma 15] D. P. Kingma, J. Ba, "Adam: A Method for Stochastic Optimization", in *International Conference on Learning Representations (ICLR)*. 2015.

[Kumar 21] A. Kumar, Z. Fu, D. Pathak, J. Malik, "RMA: Rapid Motor Adaptation for Legged Robots", in *Proceedings of Robotics: Science and Systems (RSS)*. 2021.

[Lee 20] J. Lee, J. Hwangbo, L. Wellhausen, V. Koltun, M. Hutter, "Learning Quadrupedal Locomotion over Challenging Terrain", *Science Robotics*, vol. 5, no. 47, p. eabc5986, 2020.

[Liu 25] X. Liu, C. Lu, H. Wu, B. Hu, X. Li, Z. Li, X. Guo, "FRMA: Four-Phase Rapid Motor Adaptation Framework", *Machines*, vol. 13, no. 10, p. 885, 2025.

[Ljung 99] L. Ljung, *System Identification: Theory for the User*, 2nd ed., Upper Saddle River, NJ: Prentice Hall, 1999.

[Lutter 19] M. Lutter, C. Ritter, J. Peters, "Deep Lagrangian Networks: Using Physics as Model Prior for Deep Learning", in *International Conference on Learning Representations (ICLR)*. 2019.

[Maitan 26] L. Maitan, L. Toschi, C. Zanette, E. G. Vergamini, L. F. Santos, T. Boaventura, "Contraction Metric Based Safe Reinforcement Learning Force Control for a Hydraulic Actuator with Real-World Training", arXiv preprint arXiv:2602.08977, 2026.

[Makoviychuk 21] V. Makoviychuk, L. Wawrzyniak, Y. Guo, M. Lu, K. Storey, M. Macklin, D. Hoeller, N. Rudin, A. Allshire, A. Handa, G. State, "Isaac Gym: High Performance GPU-Based Physics Simulation for Robot Learning", arXiv preprint arXiv:2108.10470, 2021.

[Mattila 17] J. Mattila, J. Koivumäki, D. G. Caldwell, C. Semini, "A Survey on Control of Hydraulic Robotic Manipulators With Projection to Future Trends", *IEEE/ASME Transactions on Mechatronics*, vol. 22, no. 2, pp. 669–680, 2017.

[McKay 79] M. D. McKay, R. J. Beckman, W. J. Conover, "A Comparison of Three Methods for Selecting Values of Input Variables in the Analysis of Output from a Computer Code", *Technometrics*, vol. 21, no. 2, pp. 239–245, 1979.

[Merritt 67] H. E. Merritt, *Hydraulic Control Systems*, New York: John Wiley & Sons, 1967.

[Mittal 23] M. Mittal, C. Yu, Q. Yu, J. Liu, N. Rudin, D. Hoeller, J. L. Yuan, R. Singh, Y. Guo, H. Mazhar, A. Mandlekar, B. Babich, G. State, M. Hutter, A. Garg, "Orbit: A Unified Simulation Framework for Interactive Robot Learning Environments", *IEEE Robotics and Automation Letters*, vol. 8, no. 6, pp. 3740–3747, 2023.

[MULAG Fahrzeugwerk 26] MULAG Fahrzeugwerk, "FME 500 Front Mower — Product Documentation", Heinz Wössner GmbH u. Co. KG, Oppenau, Germany. `https://www.mulag.de/en/roadside-equipment/products/unimog-equipment/front-mower/fme-500/`, 2026. Accessed July 2026.

[Nagabandi 18] A. Nagabandi, G. Kahn, R. S. Fearing, S. Levine, "Neural Network Dynamics for Model-Based Deep Reinforcement Learning with Model-Free Fine-Tuning", in *Proceedings of the IEEE International Conference on Robotics and Automation (ICRA)*. 2018, pp. 7559–7566.

[Nakamura 86] Y. Nakamura, H. Hanafusa, "Inverse Kinematic Solutions With Singularity Robustness for Robot Manipulator Control", *ASME Journal of Dynamic Systems, Measurement, and Control*, vol. 108, no. 3, pp. 163–171, 1986.

[Nguyen-Tuong 11] D. Nguyen-Tuong, J. Peters, "Model Learning for Robot Control: A Survey", *Cognitive Processing*, vol. 12, no. 4, pp. 319–340, 2011.

[Paszke 19] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein, L. Antiga, A. Desmaison, A. Köpf, E. Yang, Z. DeVito, M. Raison, A. Tejani, S. Chilamkurthy, B. Steiner, L. Fang, J. Bai, S. Chintala, "PyTorch: An Imperative Style, High-Performance Deep Learning Library", in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32. 2019, pp. 8024–8035.

[Peng 18] X. B. Peng, M. Andrychowicz, W. Zaremba, P. Abbeel, "Sim-to-Real Transfer of Robotic Control with Dynamics Randomization", in *Proceedings of the IEEE International Conference on Robotics and Automation (ICRA)*. 2018, pp. 3803–3810.

[Rajeswaran 18] A. Rajeswaran, V. Kumar, A. Gupta, G. Vezzani, J. Schulman, E. Todorov, S. Levine, "Learning Complex Dexterous Manipulation with Deep Reinforcement Learning and Demonstrations", in *Proceedings of Robotics: Science and Systems (RSS)*. 2018.

[Ross 11] S. Ross, G. J. Gordon, J. A. Bagnell, "A Reduction of Imitation Learning and Structured Prediction to No-Regret Online Learning", in *Proceedings of the 14th International Conference on Artificial Intelligence and Statistics (AISTATS)*. 2011, pp. 627–635.

[Rudin 22] N. Rudin, D. Hoeller, P. Reist, M. Hutter, "Learning to Walk in Minutes Using Massively Parallel Deep Reinforcement Learning", in *Proceedings of the Conference on Robot Learning (CoRL)*. 2022, pp. 91–100.

[Schoukens 19] J. Schoukens, L. Ljung, "Nonlinear System Identification: A User-Oriented Road Map", *IEEE Control Systems Magazine*, vol. 39, no. 6, pp. 28–99, 2019.

[Schulman 15] J. Schulman, S. Levine, P. Moritz, M. I. Jordan, P. Abbeel, "Trust Region Policy Optimization", in *Proceedings of the 32nd International Conference on Machine Learning (ICML)*, ser. Proceedings of Machine Learning Research, vol. 37. 2015, pp. 1889–1897.

[Schulman 16] J. Schulman, P. Moritz, S. Levine, M. Jordan, P. Abbeel, "High-Dimensional Continuous Control Using Generalized Advantage Estimation", in *International Conference on Learning Representations (ICLR)*. 2016.

[Schulman 17] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, O. Klimov, "Proximal Policy Optimization Algorithms", arXiv preprint arXiv:1707.06347, 2017.

[Siciliano 09] B. Siciliano, L. Sciavicco, L. Villani, G. Oriolo, *Robotics: Modelling, Planning and Control*, London: Springer, 2009.

[Silver 18] T. Silver, K. Allen, J. Tenenbaum, L. Kaelbling, "Residual Policy Learning", arXiv preprint arXiv:1812.06298, 2018.

[Slotine 91] J.-J. E. Slotine, W. Li, *Applied Nonlinear Control*, Englewood Cliffs, NJ: Prentice Hall, 1991.

[Sobol' 67] I. M. Sobol', "On the Distribution of Points in a Cube and the Approximate Evaluation of Integrals", *USSR Computational Mathematics and Mathematical Physics*, vol. 7, no. 4, pp. 86–112, 1967.

[Sofrony 24] J. Sofrony, M. C. Turner, C. M. Richards, "Anti-windup compensation for model reference adaptive control schemes", *International Journal of Robust and Nonlinear Control*, vol. 34, no. 15, pp. 10176–10193, 2024.

[Spinelli 24] F. A. Spinelli, P. Egli, J. Nubert, F. Nan, T. Bleumer, P. Goegler, S. Brockes, F. Hofmann, M. Hutter, "Reinforcement Learning Control for Autonomous Hydraulic Material Handling Machines with Underactuated Tools", in *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. 2024. arXiv:2410.05093.

[Srivastava 14] N. Srivastava, G. Hinton, A. Krizhevsky, I. Sutskever, R. Salakhutdinov, "Dropout: A Simple Way to Prevent Neural Networks from Overfitting", *Journal of Machine Learning Research*, vol. 15, pp. 1929–1958, 2014.

[Sutton 18] R. S. Sutton, A. G. Barto, *Reinforcement Learning: An Introduction*, 2nd ed., Cambridge, MA: MIT Press, 2018.

[The MathWorks, Inc. 24a] The MathWorks, Inc , "Simscape Fluids: User's Guide", Natick, MA, USA, 2024. `https://www.mathworks.com/products/simscape-fluids.html`.

[The MathWorks, Inc. 24b] The MathWorks, Inc , "Simscape Multibody: User's Guide", Natick, MA, USA, 2024. `https://www.mathworks.com/products/simscape-multibody.html`.

[Tobin 17] J. Tobin, R. Fong, A. Ray, J. Schneider, W. Zaremba, P. Abbeel, "Domain Randomization for Transferring Deep Neural Networks from Simulation to the Real World", in *Proceedings of the IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. 2017, pp. 23–30.

[von Szadkowski 20] K. von Szadkowski, S. Reichel, "Phobos: A Tool for Creating Complex Robot Models", *Journal of Open Source Software*, vol. 5, no. 45, p. 1326, 2020.

[Wampler 86] C. W. Wampler, "Manipulator Inverse Kinematic Solutions Based on Vector Formulations and Damped Least-Squares Methods", *IEEE Transactions on Systems, Man, and Cybernetics*, vol. 16, no. 1, pp. 93–101, 1986.

[Yamamori 26] S. Yamamori, K. Ishihara, K. Minamikawa, R. Ohmori, T. Yasaki, N. Sugimoto, J. Morimoto, "Actuator Reality Shaping for Zero-Shot Sim-to-Real Robot Learning", arXiv preprint arXiv:2607.02205, 2026.

[Yang 26] N. Yang, C. Ji, J. Chen, H. Zheng, "State-Constrained Control for Hydraulic Manipulator Position Servo Systems with Valve Dead-Band Compensation", *Actuators*, vol. 15, no. 4, p. 196, 2026.

[Yin 25] P. Yin, T. Westenbroek, S. Bagaria, K. Huang, C.-A. Cheng, A. Kolobov, A. Gupta, "Rapidly Adapting Policies to the Real World via Simulation-Guided Fine-Tuning", arXiv preprint arXiv:2502.02705, 2025.

[Yu 26] S.-Y. Yu, J.-W. Wang, M. Horii, M. Tomizuka, V. Garg, "Learning Context-Aware Neural ODE Dynamics for Adaptive Robotic Control", arXiv preprint arXiv:2606.15469, 2026.